

Lab 3 - Speech Processing

Sanjushree Rajan

BL.EN.U4AIE23130

Objective 1

To identify, segment, and visualize different phoneme classes in a recorded speech signal and to analyze voiced and unvoiced regions using waveform-level inspection.

Record the sentence: "*She judged the blue pen, put it on the table, and quietly whispered yes.*"

Using the procedures implemented in Lab-2, perform the following steps:

1. Load and preprocess the recorded speech signal.
2. Automatically segment the speech into phoneme-level units (similar to previous lab).
3. Plot the time-domain waveform corresponding to each extracted phoneme.
4. For every phoneme segment:
 - Identify its phoneme category (vowel, plosive, fricative, affricate, semivowel, diphthong, or whisper segment).
 - Label the phoneme accordingly.
5. Select and extract:
 - One voiced phoneme segment
 - One unvoiced phoneme segment
6. Compare the waveform characteristics of the voiced and unvoiced segments and explain the observed differences based on:
 - Periodicity
 - Amplitude patterns
 - Presence or absence of noise-like structure

1. Library Imports

```
In [1]: import torch
import torchaudio
import numpy as np
import matplotlib.pyplot as plt
import librosa

from transformers import Wav2Vec2Processor, Wav2Vec2ForCTC

import warnings
warnings.filterwarnings("ignore")
```

```
C:\Users\Srijjay\AppData\Roaming\Python\Python313\site-packages\tqdm\auto.py:21:  
TqdmWarning: IPython not found. Please update jupyter and ipywidgets. See http  
s://ipywidgets.readthedocs.io/en/stable/user_install.html  
from .autonotebook import tqdm as notebook_tqdm
```

2. Load Pre-trained Wav2Vec2 Model

```
In [2]: processor = Wav2Vec2Processor.from_pretrained("facebook/wav2vec2-base-960h")  
model = Wav2Vec2ForCTC.from_pretrained("facebook/wav2vec2-base-960h")  
model.eval()
```

```
Loading weights: 100%|██████████| 212/212 [00:01<00:00, 184.18it/s, Materializing  
param=wav2vec2.feature_projection.projection.weight]  
Wav2Vec2ForCTC LOAD REPORT from: facebook/wav2vec2-base-960h  
Key | Status |  
-----+-----+  
wav2vec2.masked_spec_embed | MISSING |
```

Notes:

- MISSING :those params were newly initialized because missing from the checkpoint. Consider training on your downstream task.

```
Out[2]: Wav2Vec2ForCTC(  
    (wav2vec2): Wav2Vec2Model(  
        (feature_extractor): Wav2Vec2FeatureEncoder(  
            (conv_layers): ModuleList(  
                (0): Wav2Vec2GroupNormConvLayer(  
                    (conv): Conv1d(1, 512, kernel_size=(10,), stride=(5,), bias=False)  
                    (activation): GELUActivation()  
                    (layer_norm): GroupNorm(512, 512, eps=1e-05, affine=True)  
                )  
                (1-4): 4 x Wav2Vec2NoLayerNormConvLayer(  
                    (conv): Conv1d(512, 512, kernel_size=(3,), stride=(2,), bias=False)  
                    (activation): GELUActivation()  
                )  
                (5-6): 2 x Wav2Vec2NoLayerNormConvLayer(  
                    (conv): Conv1d(512, 512, kernel_size=(2,), stride=(2,), bias=False)  
                    (activation): GELUActivation()  
                )  
            )  
        )  
        (feature_projection): Wav2Vec2FeatureProjection(  
            (layer_norm): LayerNorm((512,), eps=1e-05, elementwise_affine=True)  
            (projection): Linear(in_features=512, out_features=768, bias=True)  
            (dropout): Dropout(p=0.1, inplace=False)  
        )  
    (encoder): Wav2Vec2Encoder(  
        (pos_conv_embed): Wav2Vec2PositionalConvEmbedding(  
            (conv): ParametrizedConv1d(  
                768, 768, kernel_size=(128,), stride=(1,), padding=(64,), groups=16  
                (parametrizations): ModuleDict(  
                    (weight): ParametrizationList(  
                        (0): _WeightNorm()  
                    )  
                )  
            )  
            (padding): Wav2Vec2SamePadLayer()  
            (activation): GELUActivation()  
        )  
        (layer_norm): LayerNorm((768,), eps=1e-05, elementwise_affine=True)  
        (dropout): Dropout(p=0.1, inplace=False)  
        (layers): ModuleList(  
            (0-11): 12 x Wav2Vec2EncoderLayer(  
                (attention): Wav2Vec2Attention(  
                    (k_proj): Linear(in_features=768, out_features=768, bias=True)  
                    (v_proj): Linear(in_features=768, out_features=768, bias=True)  
                    (q_proj): Linear(in_features=768, out_features=768, bias=True)  
                    (out_proj): Linear(in_features=768, out_features=768, bias=True)  
                )  
                (dropout): Dropout(p=0.1, inplace=False)  
                (layer_norm): LayerNorm((768,), eps=1e-05, elementwise_affine=True)  
                (feed_forward): Wav2Vec2FeedForward(  
                    (intermediate_dropout): Dropout(p=0.1, inplace=False)  
                    (intermediate_dense): Linear(in_features=768, out_features=3072, bias=True)  
                )  
                (intermediate_act_fn): GELUActivation()  
                (output_dense): Linear(in_features=3072, out_features=768, bias=True)  
            )  
            (output_dropout): Dropout(p=0.1, inplace=False)  
        )  
        (final_layer_norm): LayerNorm((768,), eps=1e-05, elementwise_affine=True)  
    )
```

```

        )
    )
)
(dropout): Dropout(p=0.1, inplace=False)
(lm_head): Linear(in_features=768, out_features=32, bias=True)
)

```

3. Load and Preprocess Speech Signal

In [3]:

```

import os
import librosa
import torch

audio_filename = "sentence1.wav"

if not os.path.exists(audio_filename):
    print(f"Error: '{audio_filename}' not found. Please upload your recorded aud
else:
    data, sample_rate = librosa.load(audio_filename, sr=None, mono=False)

    waveform = torch.from_numpy(data).float()

    if len(waveform.shape) == 1:
        waveform = waveform.unsqueeze(0)

    if waveform.shape[0] > 1:
        waveform = torch.mean(waveform, dim=0, keepdim=True)

    waveform = waveform / torch.max(torch.abs(waveform))

    print("Sampling Rate:", sample_rate)

```

Sampling Rate: 48000

4. Plot Original Speech Waveform

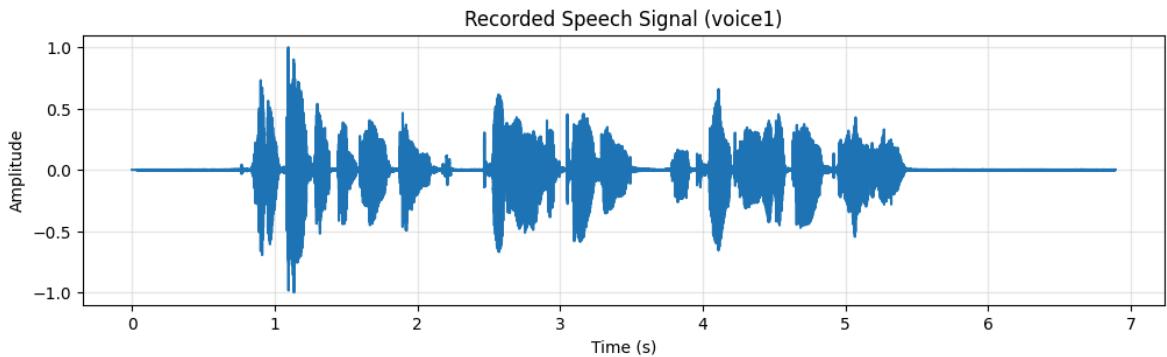
In [4]:

```

time_axis = np.arange(waveform.shape[1]) / sample_rate

plt.figure(figsize=(12,3))
plt.plot(time_axis, waveform[0].numpy())
plt.title("Recorded Speech Signal (voice1)")
plt.xlabel("Time (s)")
plt.ylabel("Amplitude")
plt.grid(alpha=0.3)
plt.show()

```



5. Recognize Phonemes using Wav2Vec2

```
In [5]: target_sampling_rate = 16000

if sample_rate != target_sampling_rate:
    waveform_np = waveform[0].numpy()
    waveform_resampled = librosa.resample(waveform_np, orig_sr=sample_rate, targ
    input_sr = target_sampling_rate
else:
    waveform_resampled = waveform[0].numpy()
    input_sr = sample_rate

input_values = processor(
    waveform_resampled,
    sampling_rate=input_sr,
    return_tensors="pt"
).input_values

with torch.no_grad():
    logits = model(input_values).logits

predicted_ids = torch.argmax(logits, dim=-1)
transcription = processor.decode(predicted_ids[0])

print("Recognized Phoneme Sequence:")
print(transcription)
```

Recognized Phoneme Sequence:

SHE JUDGED THE BLUE PEN PUT IT ON THE TABLE AND QUIETLY WHISPERED YES

6. Estimate Phoneme Time Intervals

```
In [6]: predicted_tokens = predicted_ids[0].numpy()
time_step = waveform.shape[1] / logits.shape[1] / sample_rate

phoneme_intervals = []
current_phoneme = None
start_idx = 0

for idx, token_id in enumerate(predicted_tokens):
    phoneme = processor.decode([token_id]).strip()

    if phoneme and phoneme != current_phoneme:
        if current_phoneme is not None:
            phoneme_intervals.append({
```

```

        "phoneme": current_phoneme,
        "start_sample": int(start_idx * time_step * sample_rate),
        "end_sample": int(idx * time_step * sample_rate)
    })
    current_phoneme = phoneme
    start_idx = idx

if current_phoneme is not None:
    phoneme_intervals.append({
        "phoneme": current_phoneme,
        "start_sample": int(start_idx * time_step * sample_rate),
        "end_sample": waveform.shape[1]
    })

print(f"Detected {len(phoneme_intervals)} phoneme segments")

```

Detected 56 phoneme segments

7. Phoneme Category Classification

```
In [7]: phoneme_categories = {
    'A':'vowel', 'E':'vowel', 'I':'vowel', 'O':'vowel', 'U':'vowel',
    'P':'plosive', 'B':'plosive', 'T':'plosive', 'D':'plosive', 'K':'plosive',
    'F':'fricative', 'S':'fricative', 'H':'fricative', 'Z':'fricative',
    'Y':'semivowel', 'W':'semivowel', 'L':'semivowel', 'R':'semivowel'
}

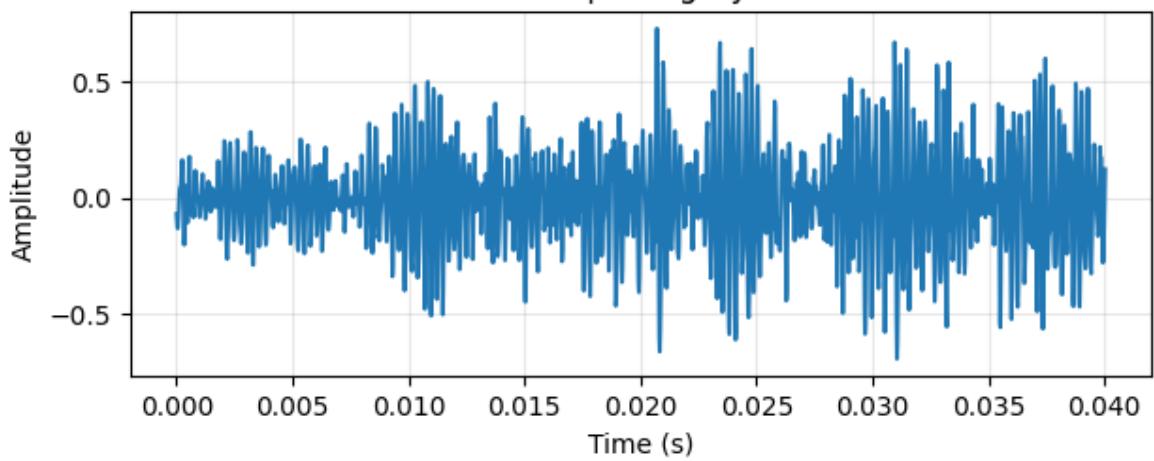
for p in phoneme_intervals:
    p["category"] = phoneme_categories.get(p["phoneme"].upper(), "whisper/other")
```

8. Extract & Plot Each Phoneme Waveform

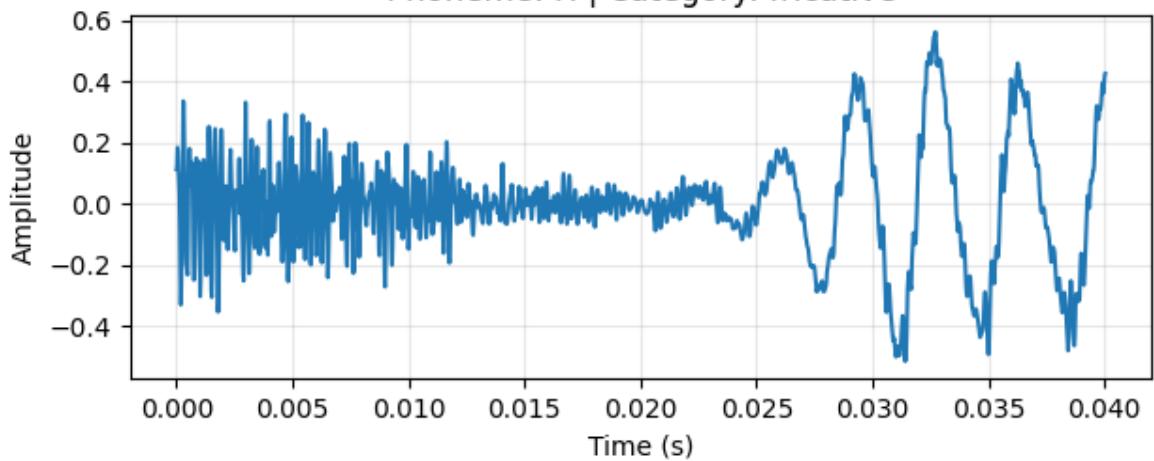
```
In [8]: def plot_phoneme(segment, sr, phoneme, category):
    t = np.arange(len(segment)) / sr
    plt.figure(figsize=(7,2.5))
    plt.plot(t, segment)
    plt.title(f"Phoneme: {phoneme} | Category: {category}")
    plt.xlabel("Time (s)")
    plt.ylabel("Amplitude")
    plt.grid(alpha=0.3)
    plt.show()

for p in phoneme_intervals:
    segment = waveform[0, p["start_sample"]:p["end_sample"]].numpy()
    plot_phoneme(segment, sample_rate, p["phoneme"], p["category"])
```

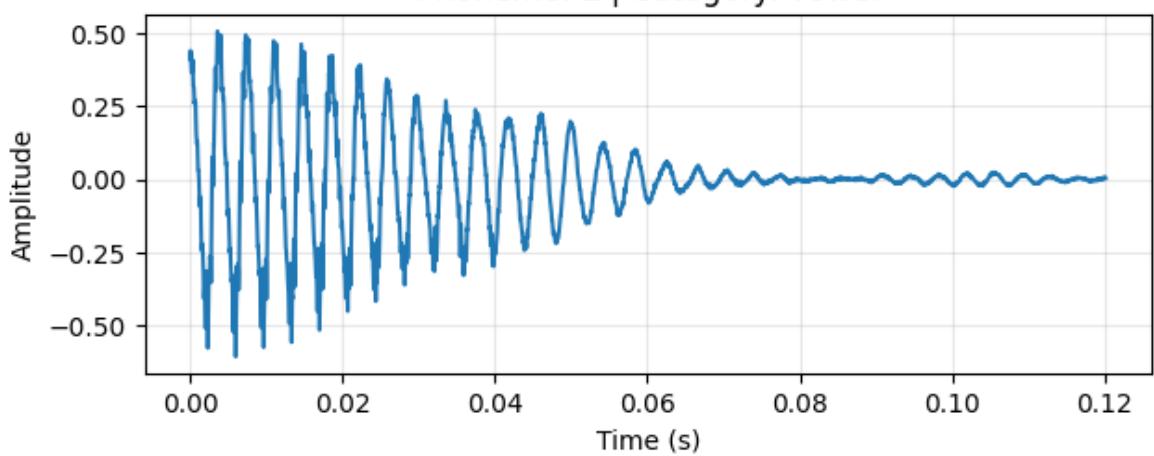
Phoneme: S | Category: fricative



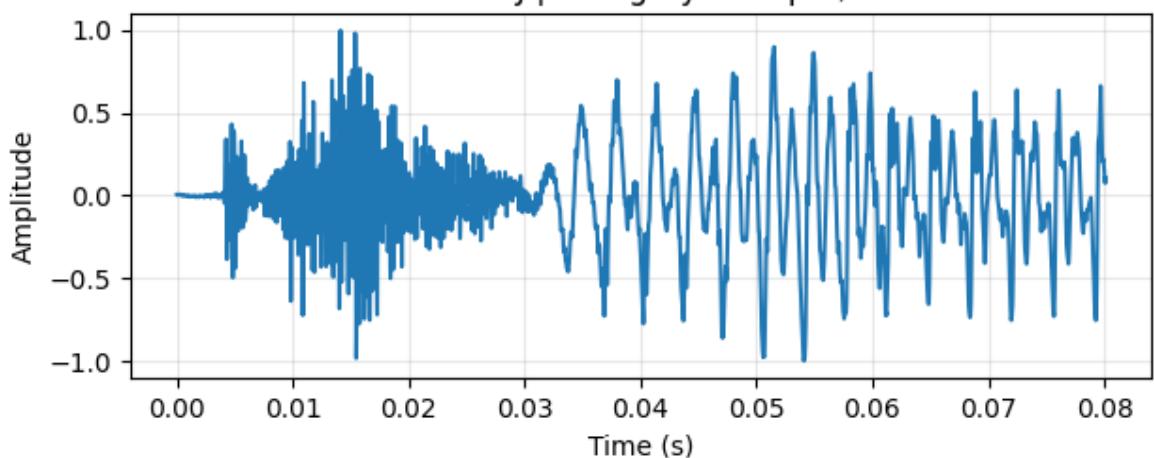
Phoneme: H | Category: fricative



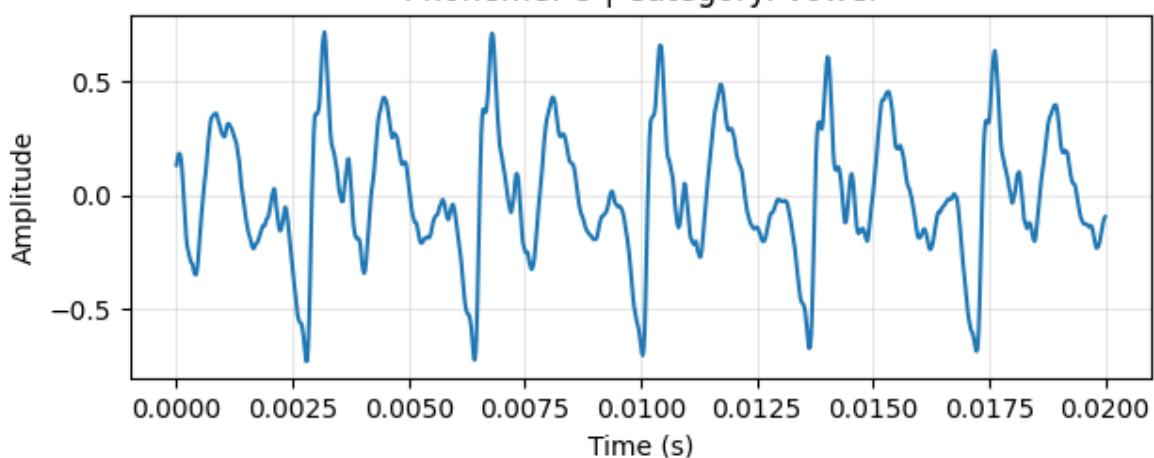
Phoneme: E | Category: vowel



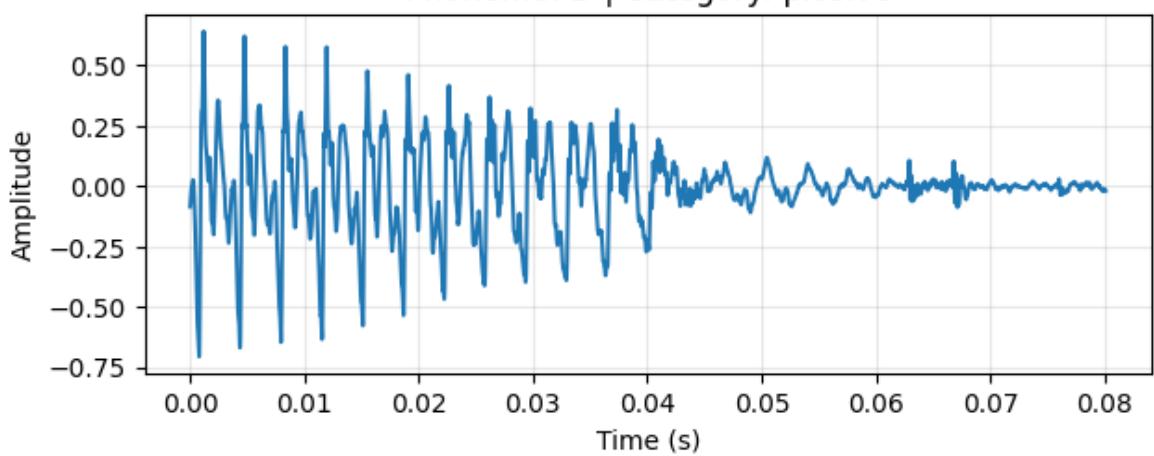
Phoneme: J | Category: whisper/other



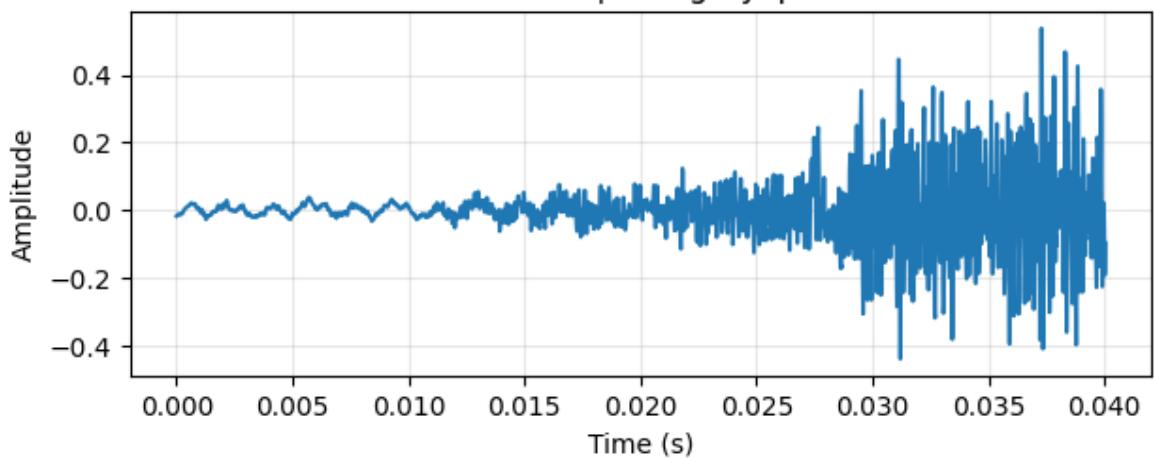
Phoneme: U | Category: vowel



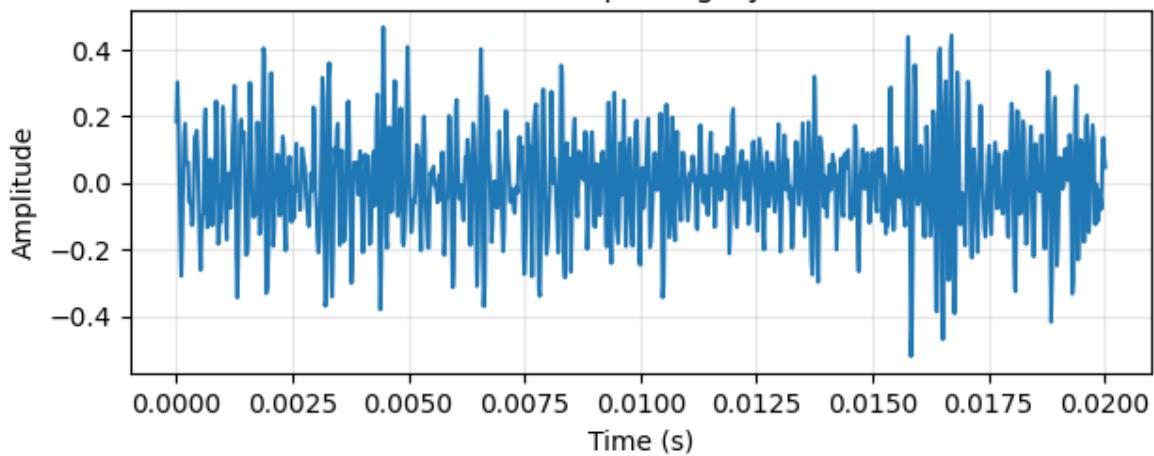
Phoneme: D | Category: plosive



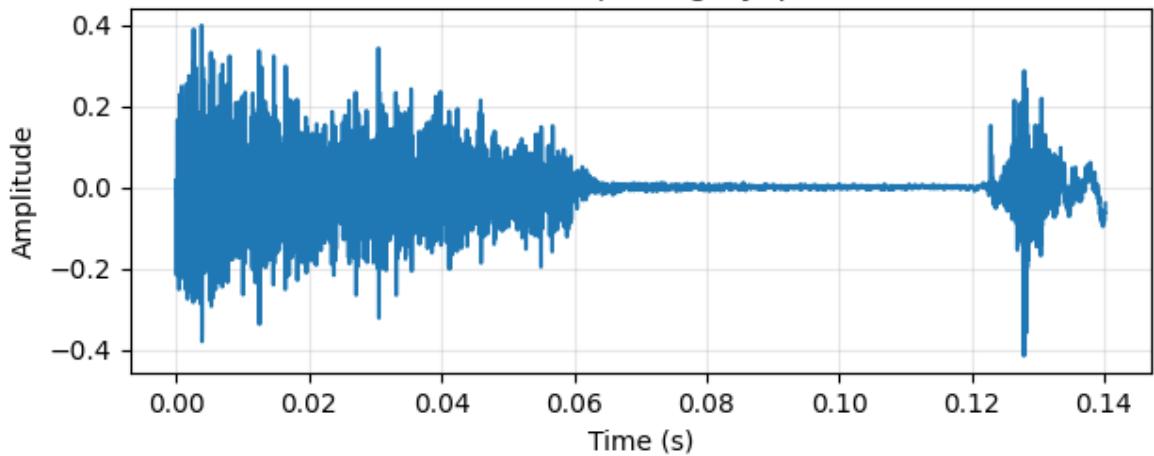
Phoneme: G | Category: plosive



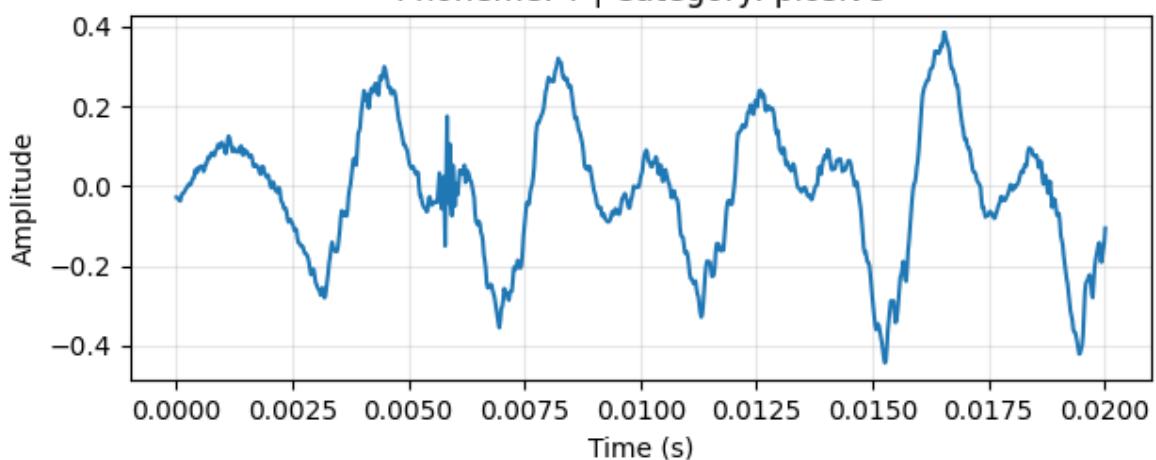
Phoneme: E | Category: vowel



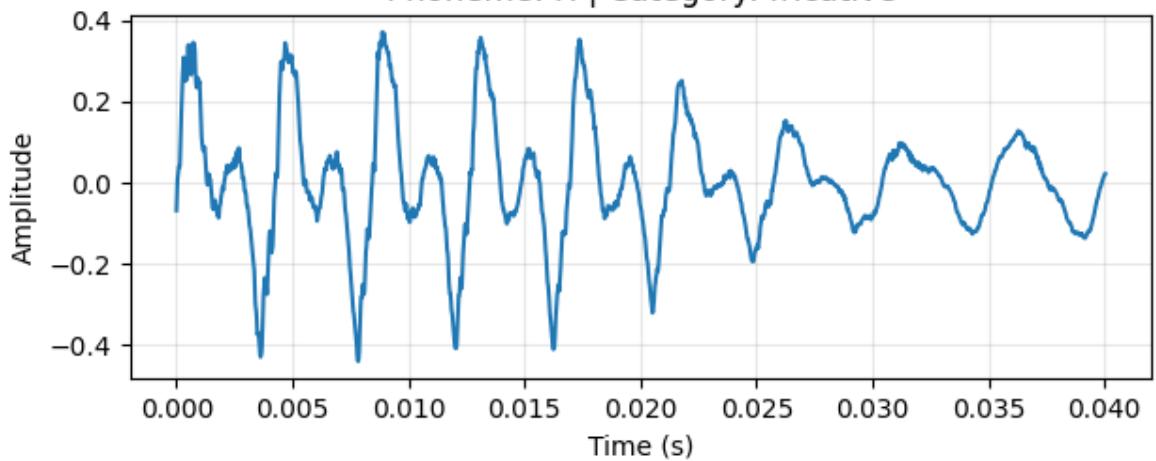
Phoneme: D | Category: plosive



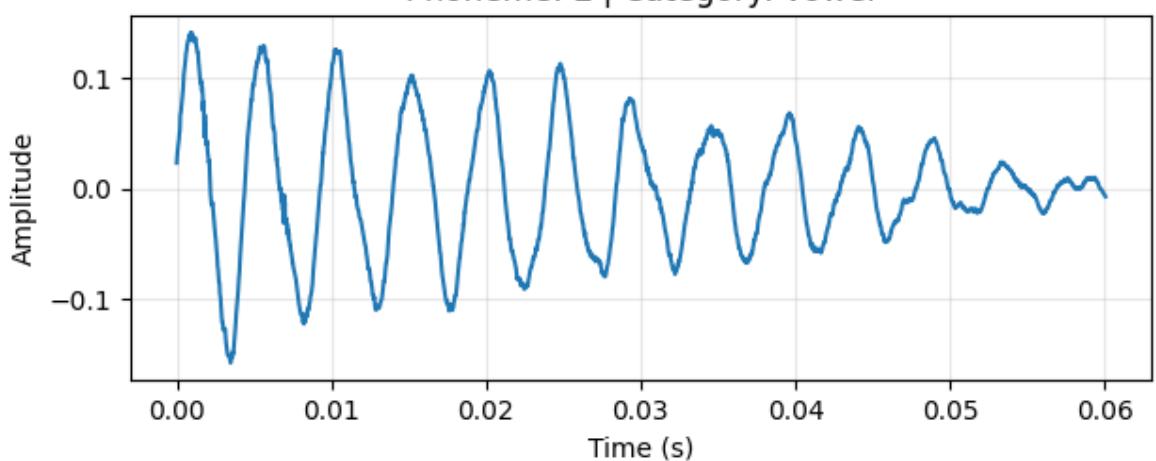
Phoneme: T | Category: plosive



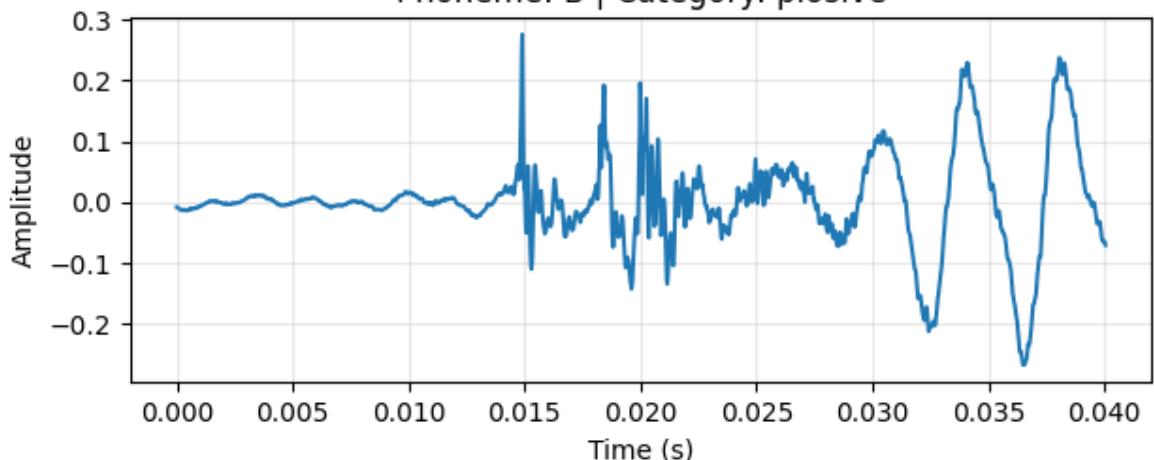
Phoneme: H | Category: fricative



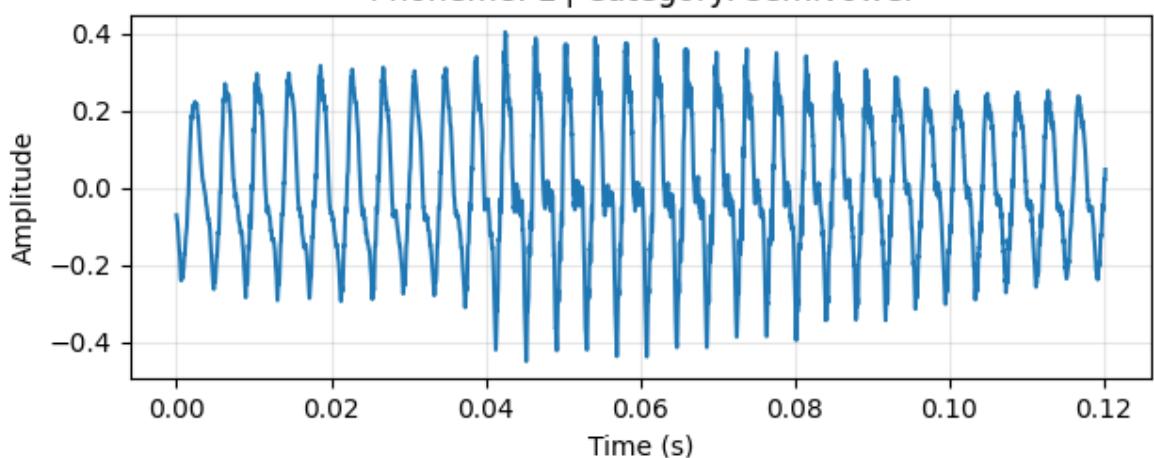
Phoneme: E | Category: vowel



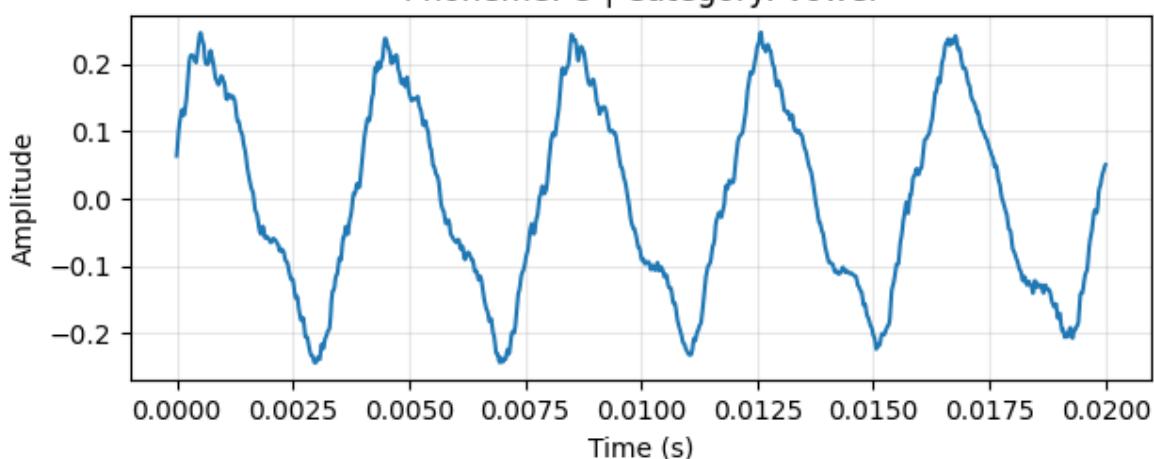
Phoneme: B | Category: plosive



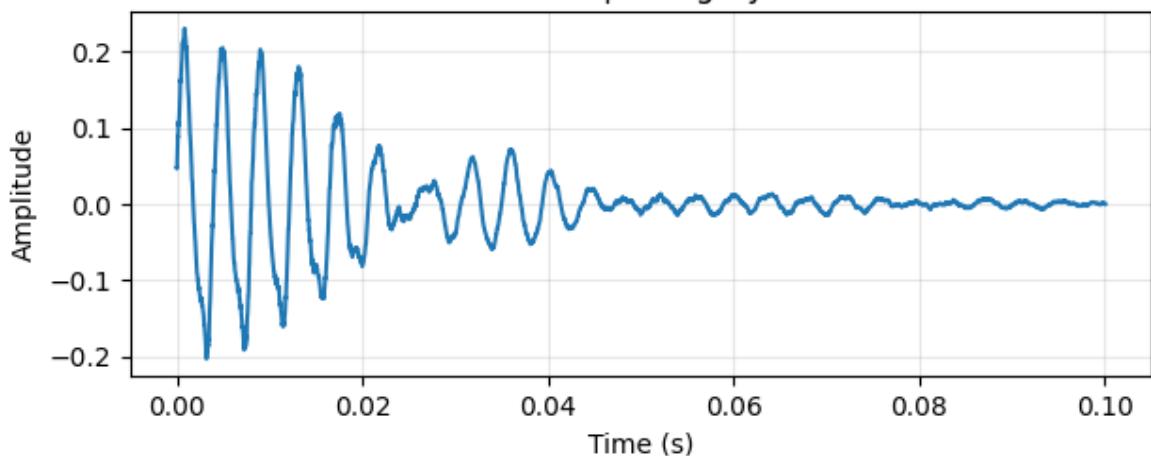
Phoneme: L | Category: semivowel



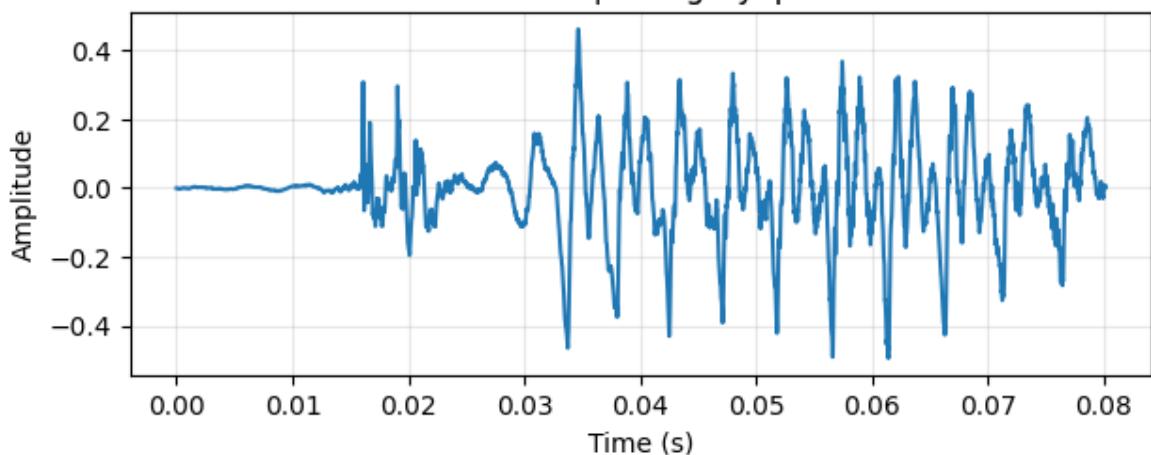
Phoneme: U | Category: vowel



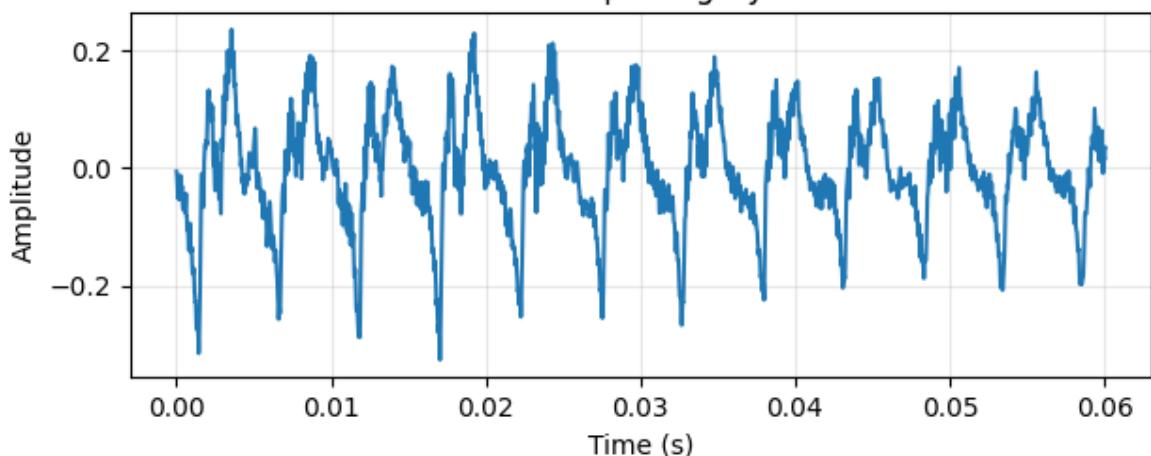
Phoneme: E | Category: vowel



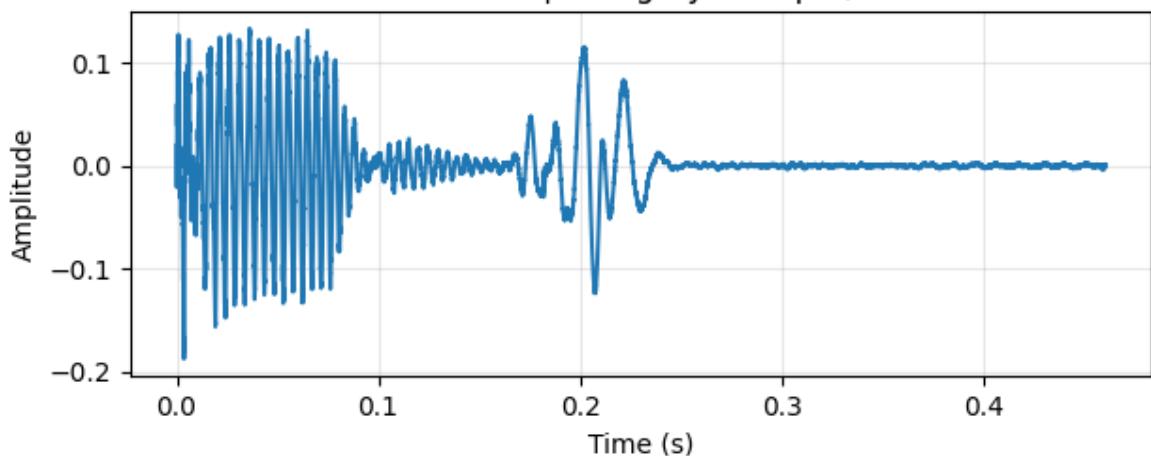
Phoneme: P | Category: plosive



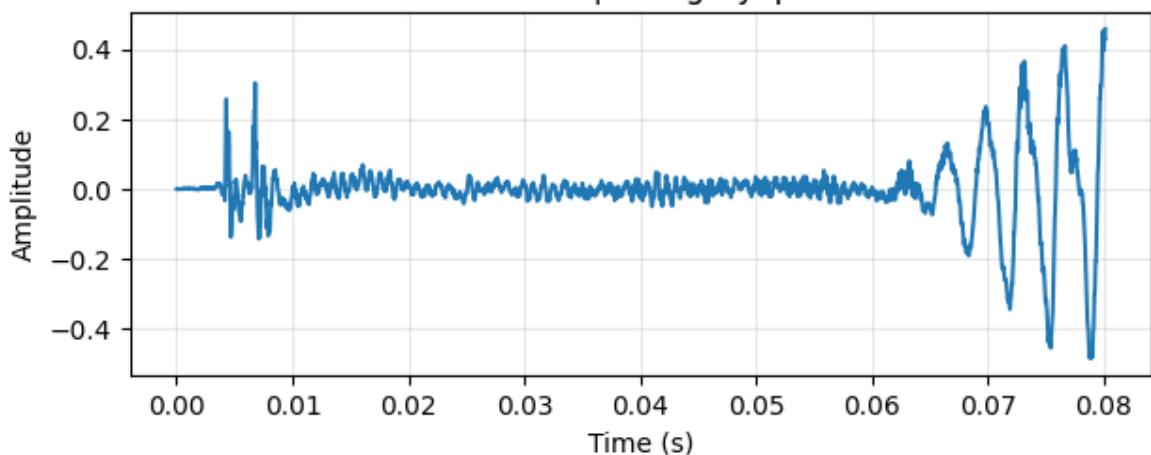
Phoneme: E | Category: vowel



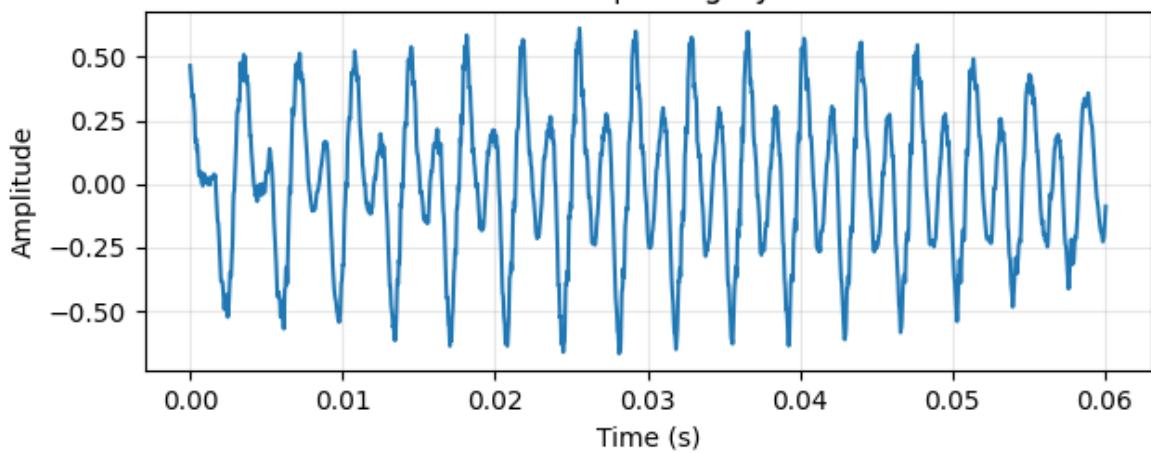
Phoneme: N | Category: whisper/other



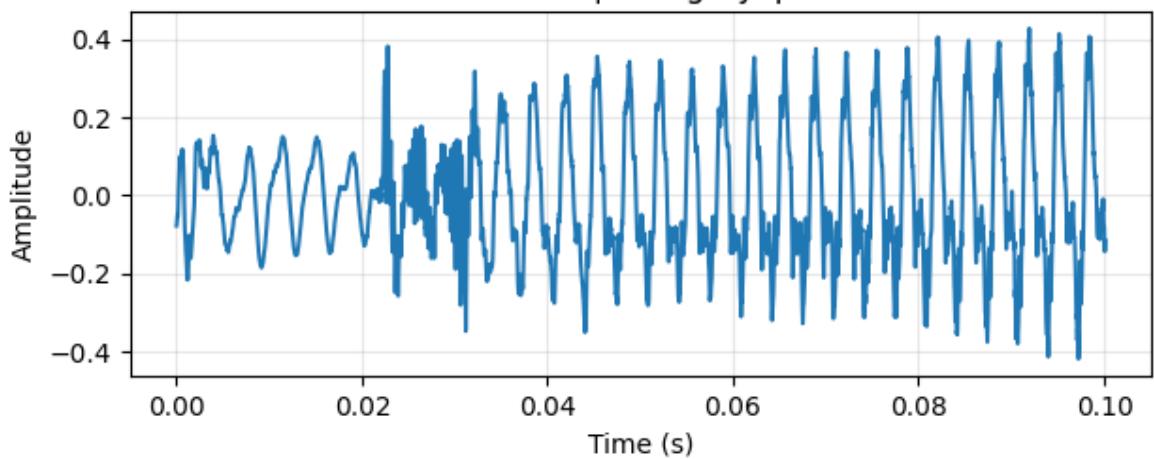
Phoneme: P | Category: plosive



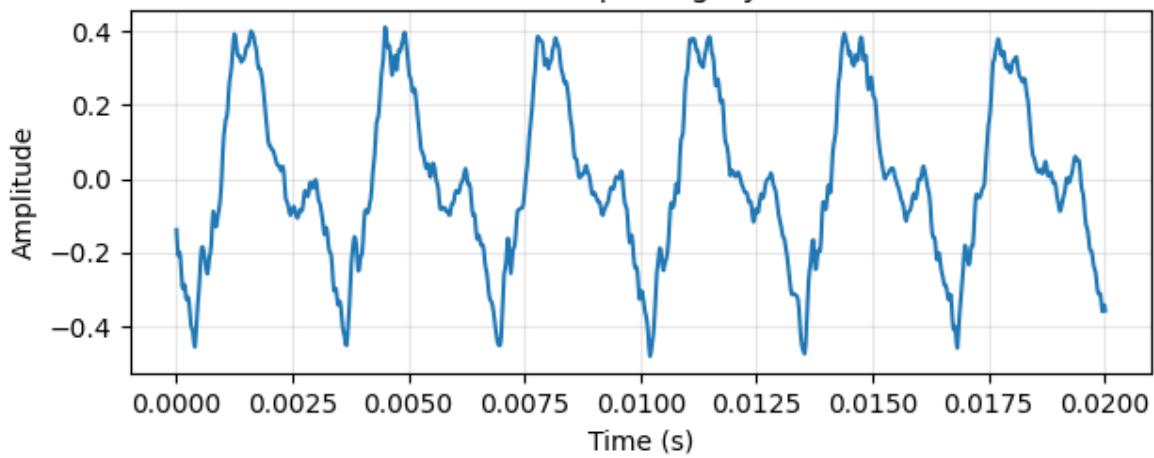
Phoneme: U | Category: vowel



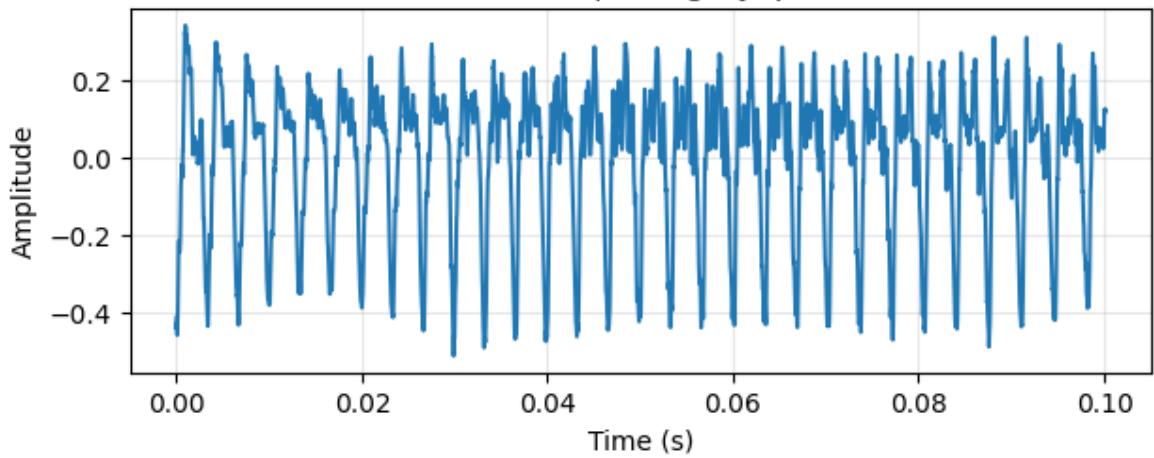
Phoneme: T | Category: plosive



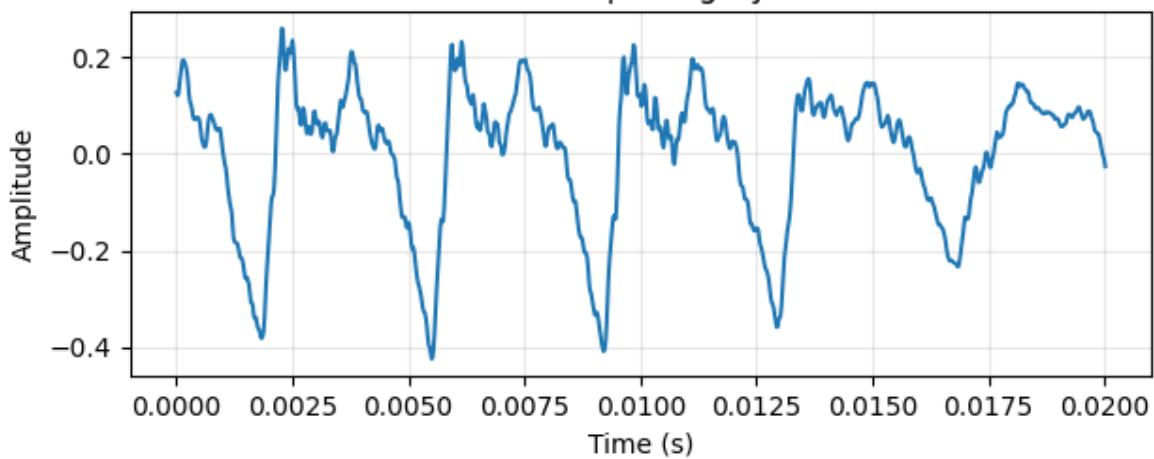
Phoneme: I | Category: vowel



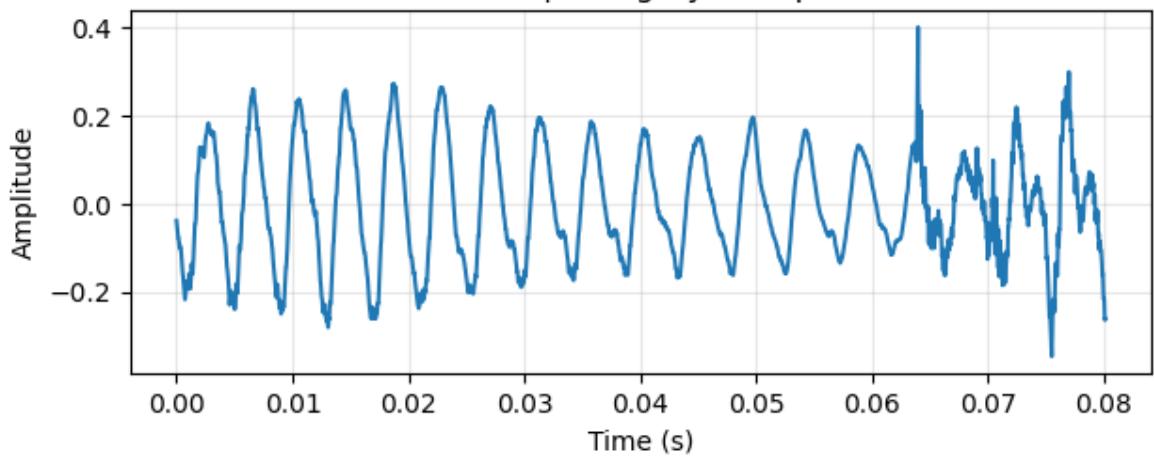
Phoneme: T | Category: plosive



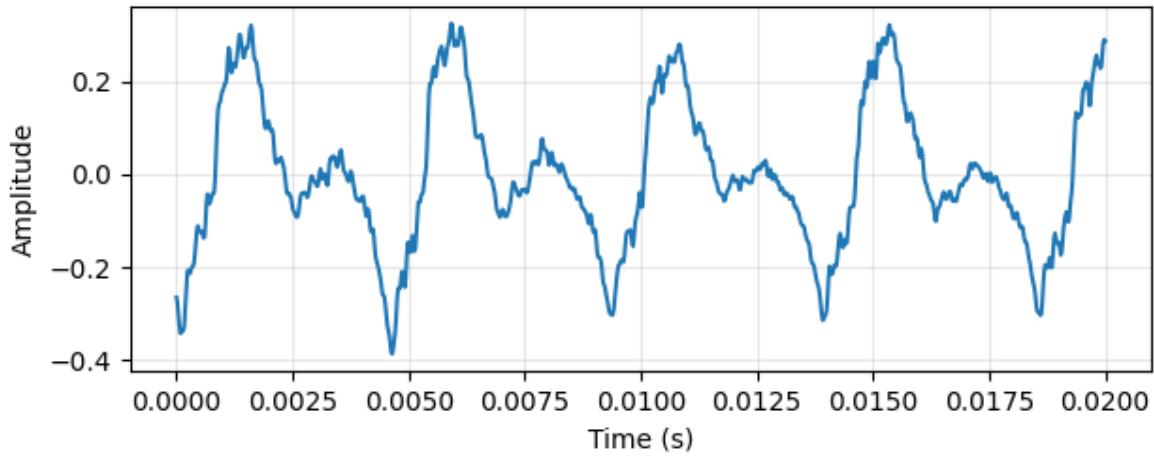
Phoneme: O | Category: vowel



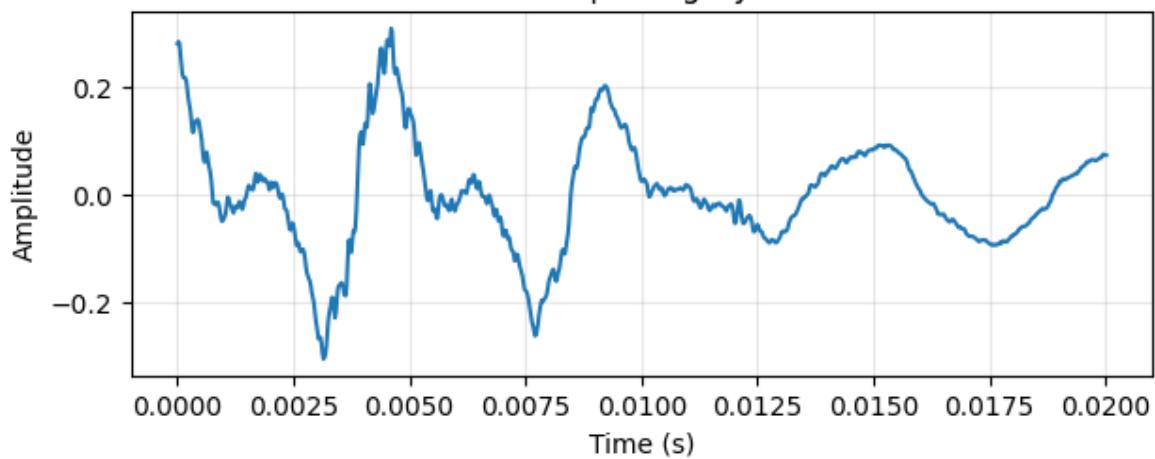
Phoneme: N | Category: whisper/other



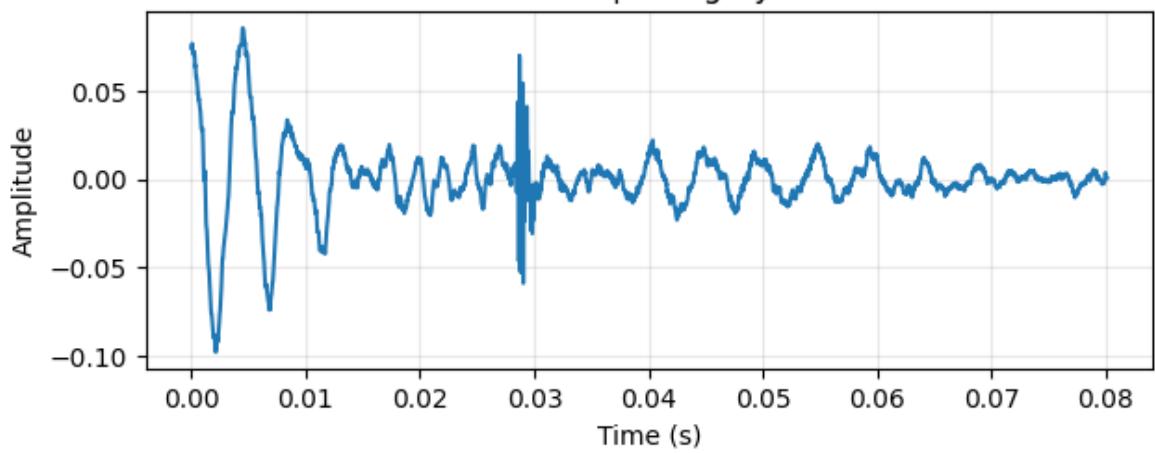
Phoneme: T | Category: plosive



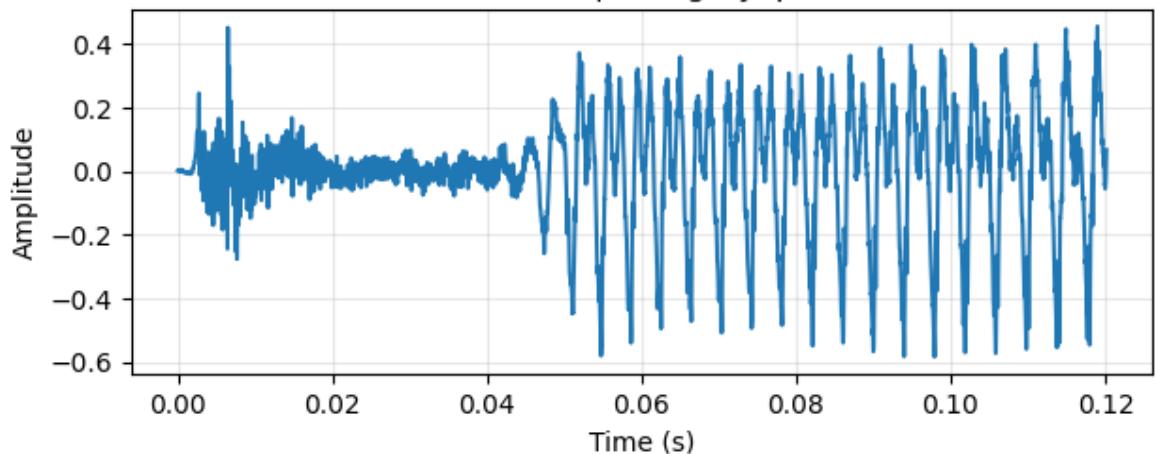
Phoneme: H | Category: fricative



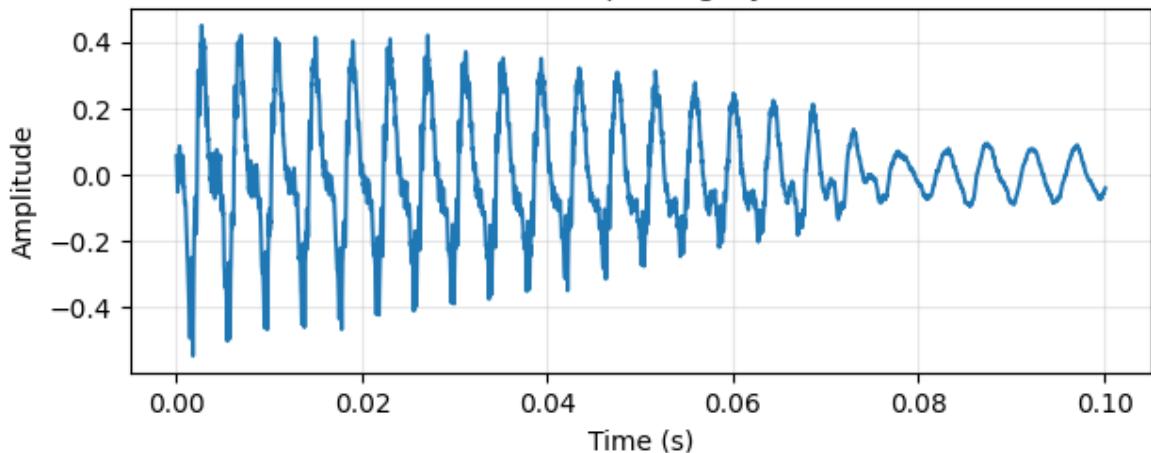
Phoneme: E | Category: vowel



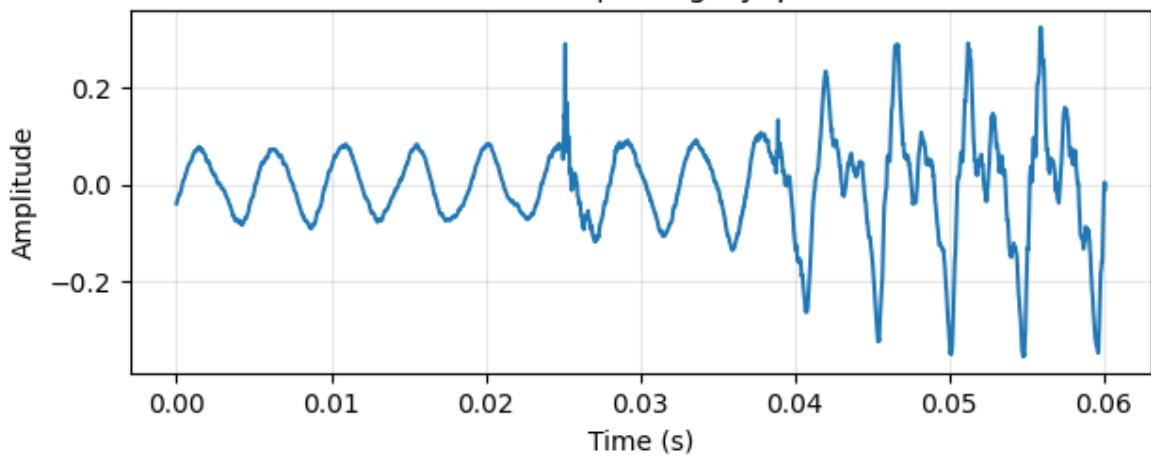
Phoneme: T | Category: plosive



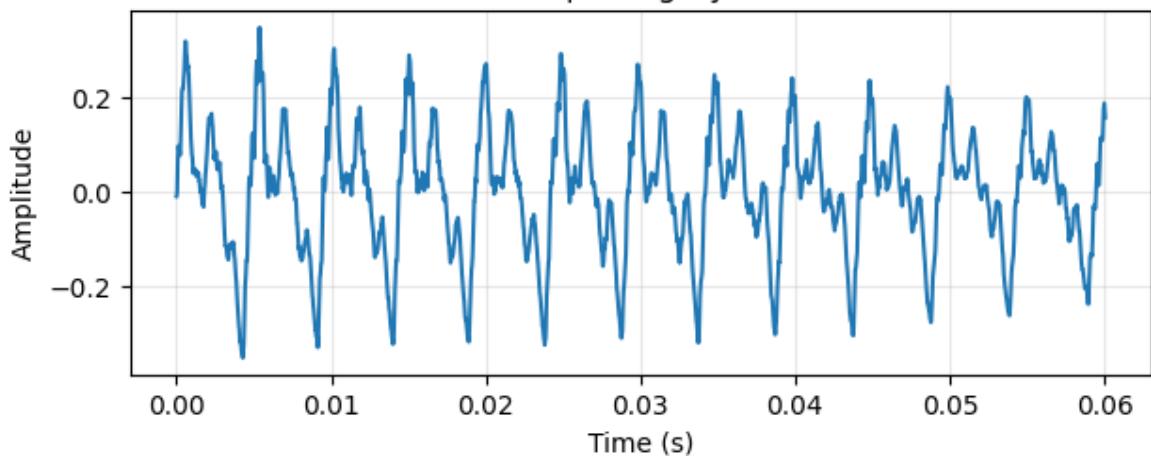
Phoneme: A | Category: vowel



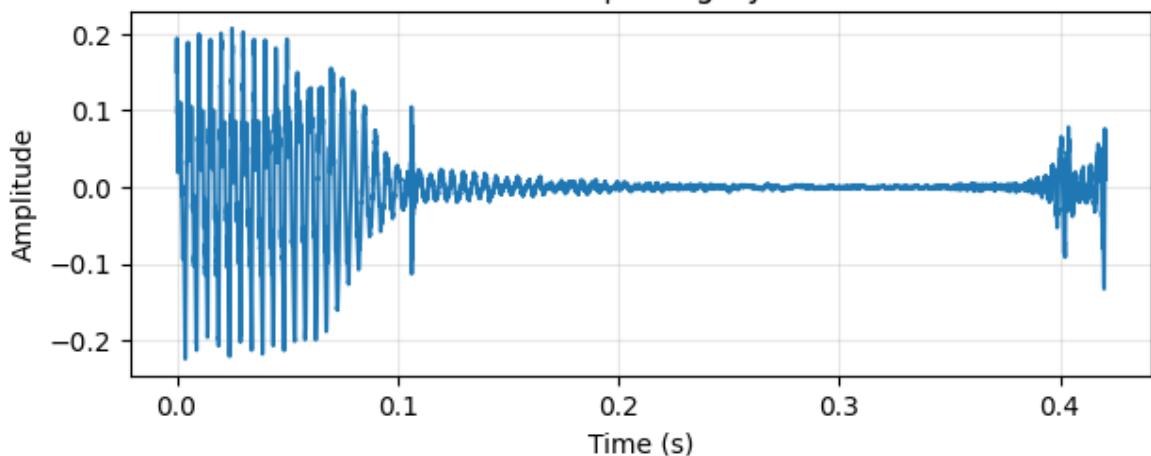
Phoneme: B | Category: plosive



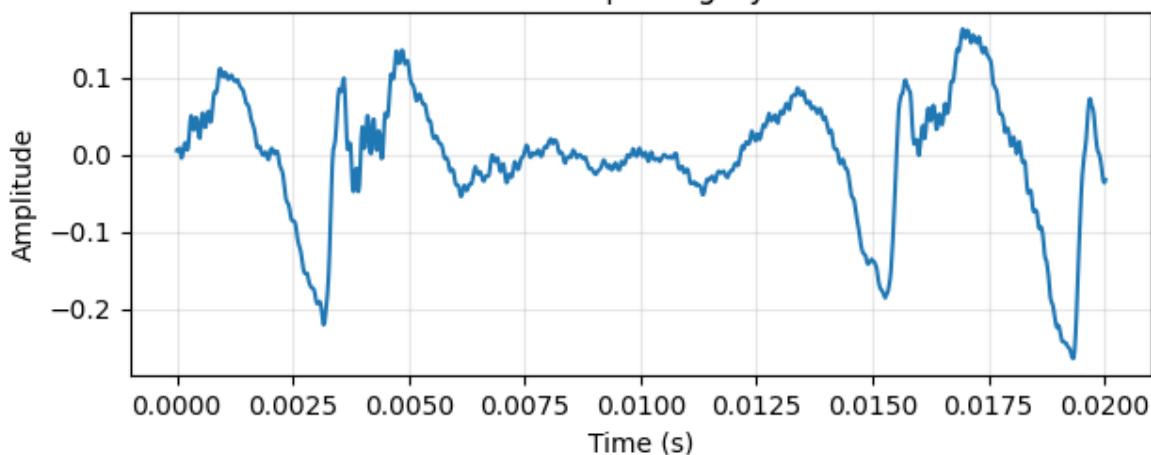
Phoneme: L | Category: semivowel



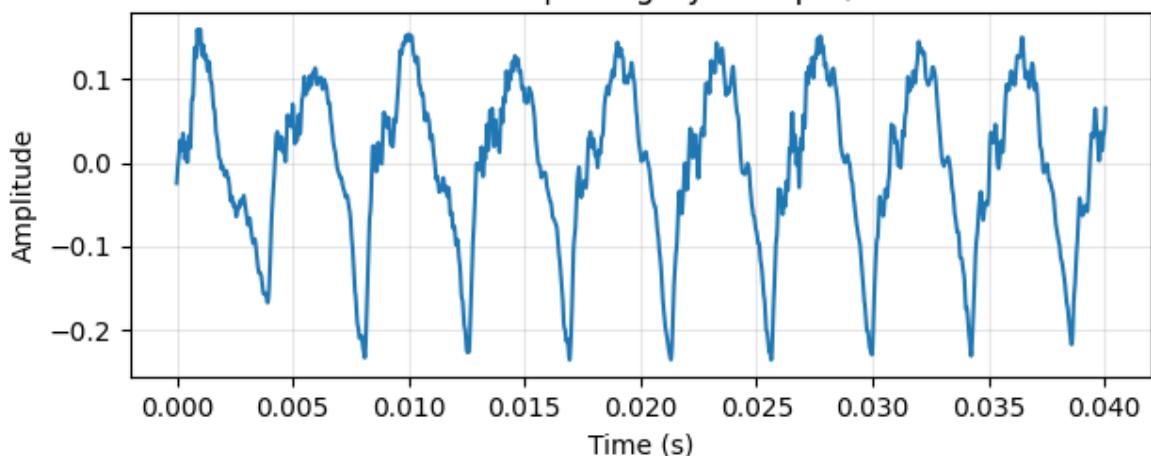
Phoneme: E | Category: vowel



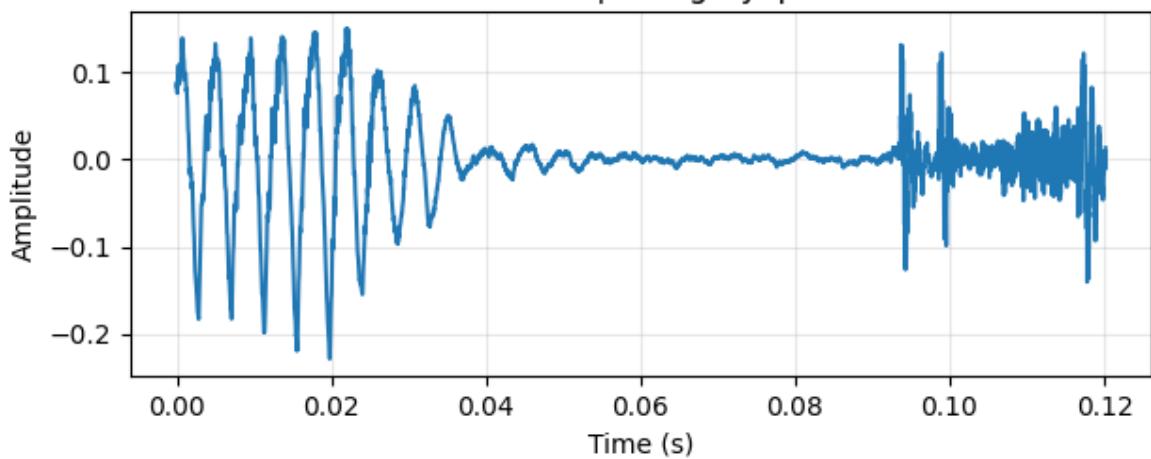
Phoneme: A | Category: vowel



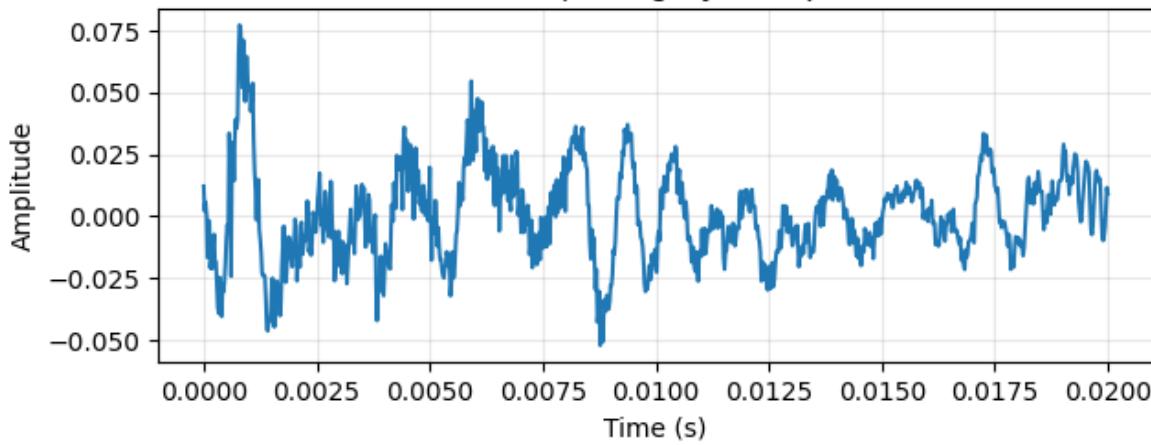
Phoneme: N | Category: whisper/other



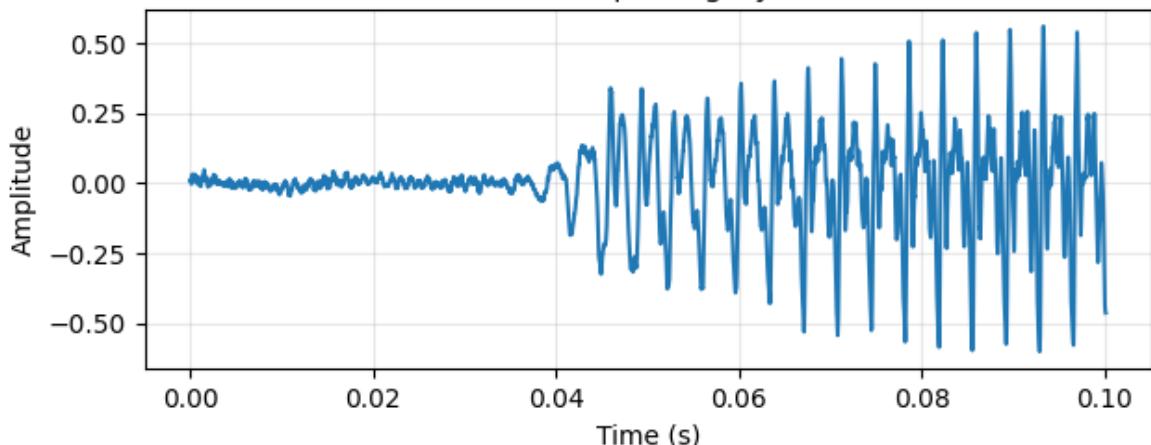
Phoneme: D | Category: plosive



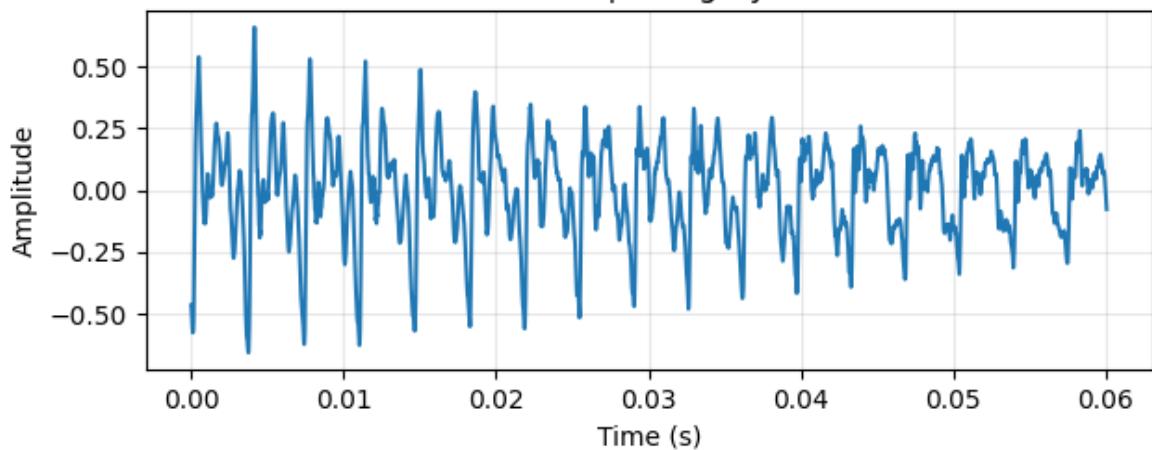
Phoneme: Q | Category: whisper/other



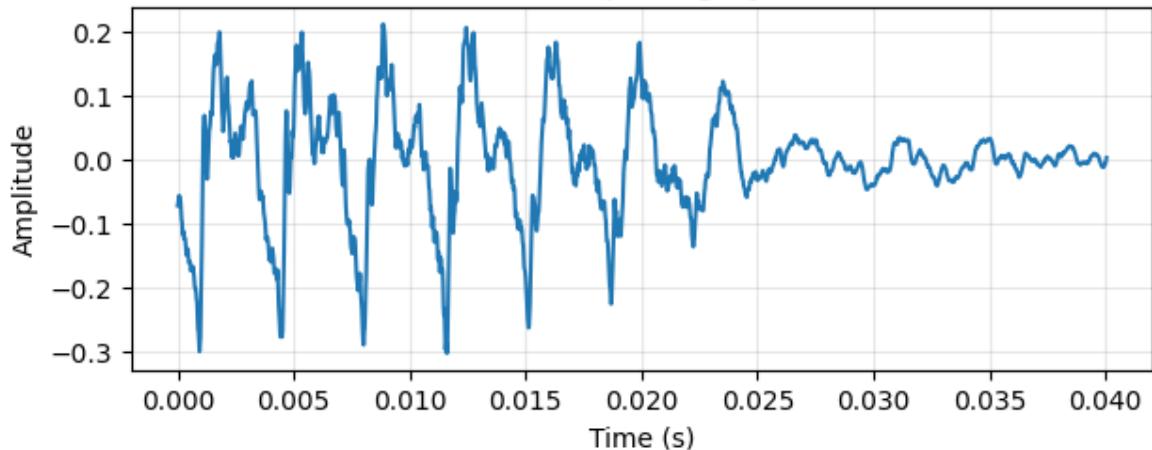
Phoneme: U | Category: vowel



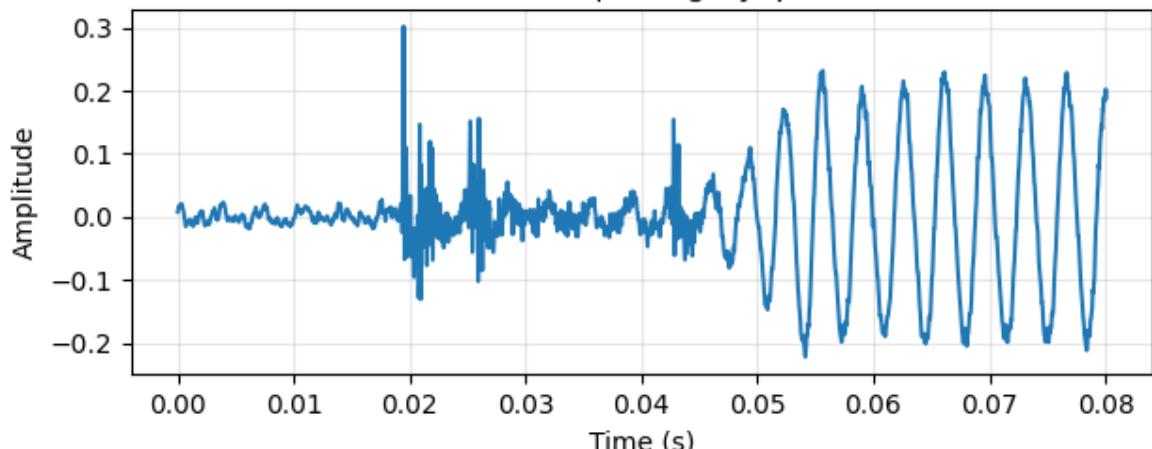
Phoneme: I | Category: vowel

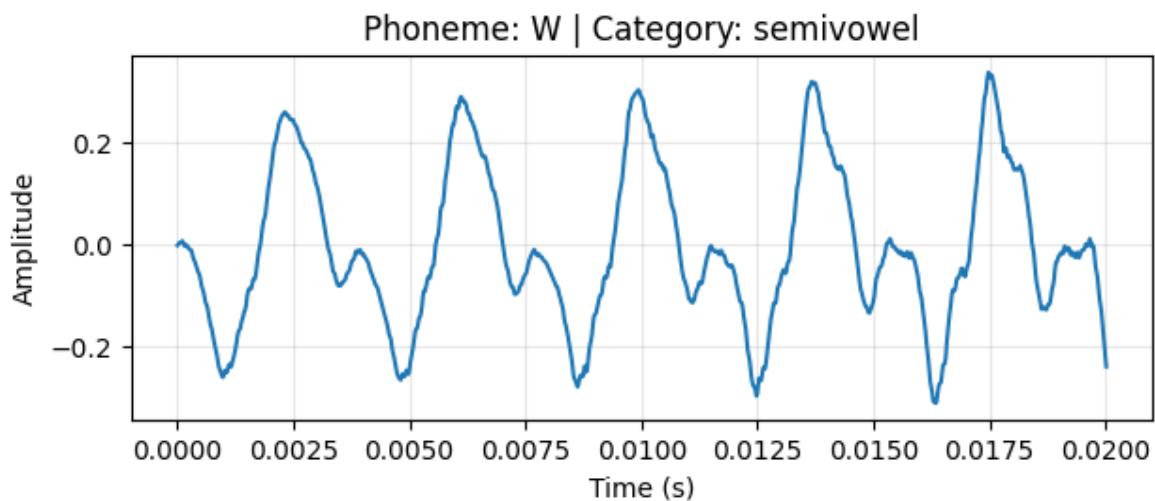
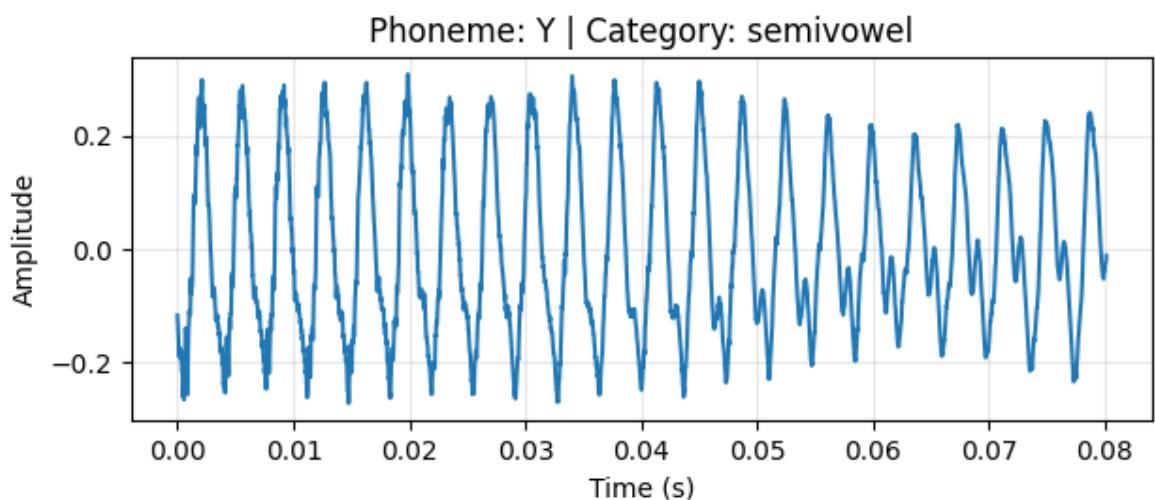
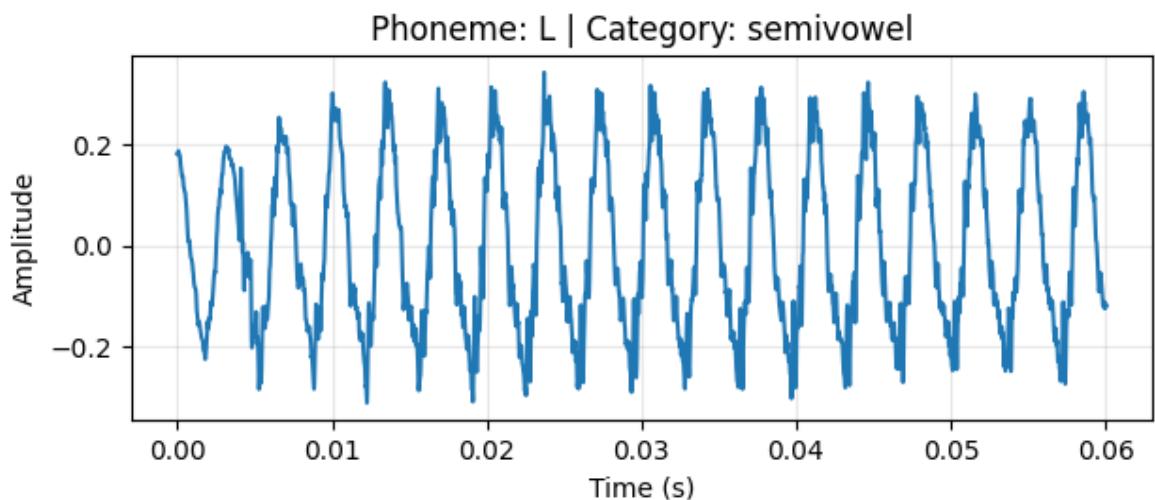


Phoneme: E | Category: vowel

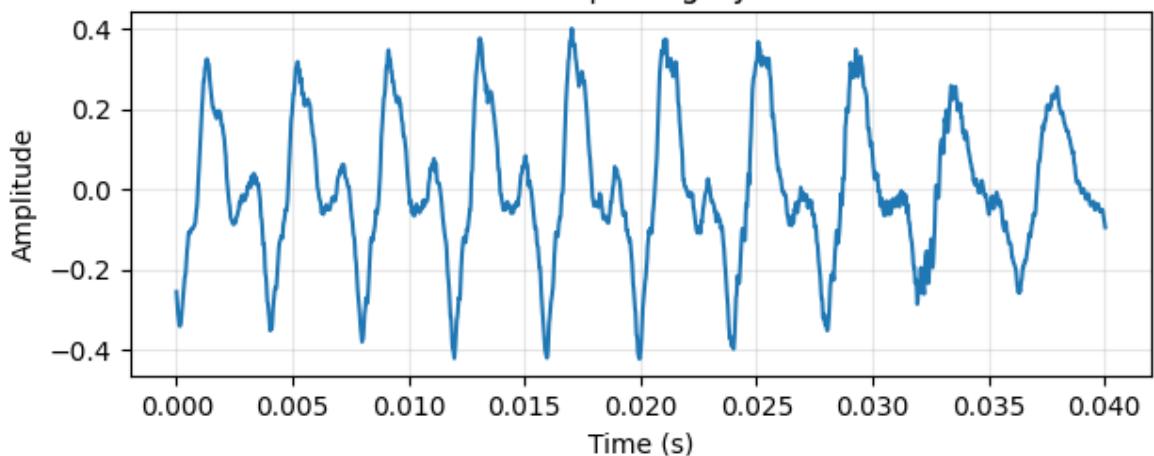


Phoneme: T | Category: plosive

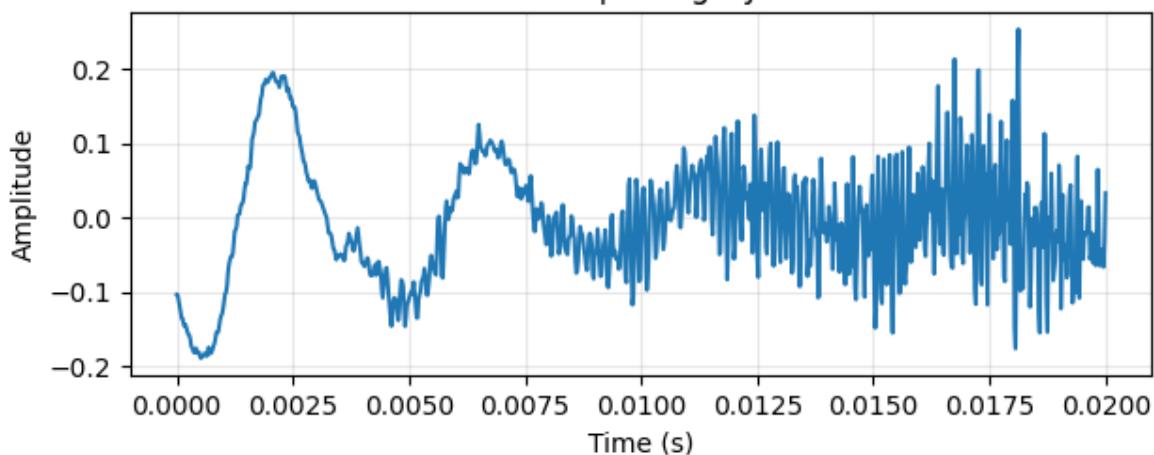




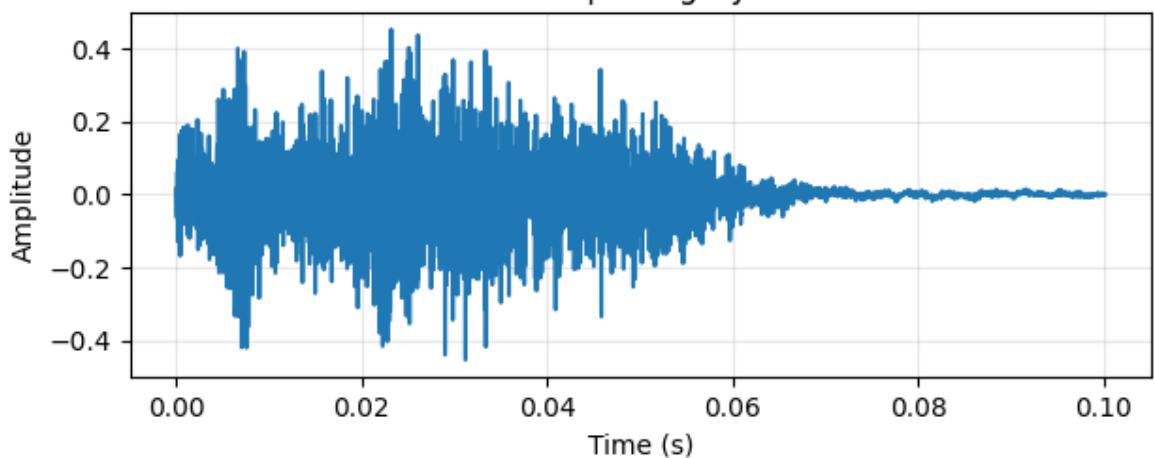
Phoneme: H | Category: fricative

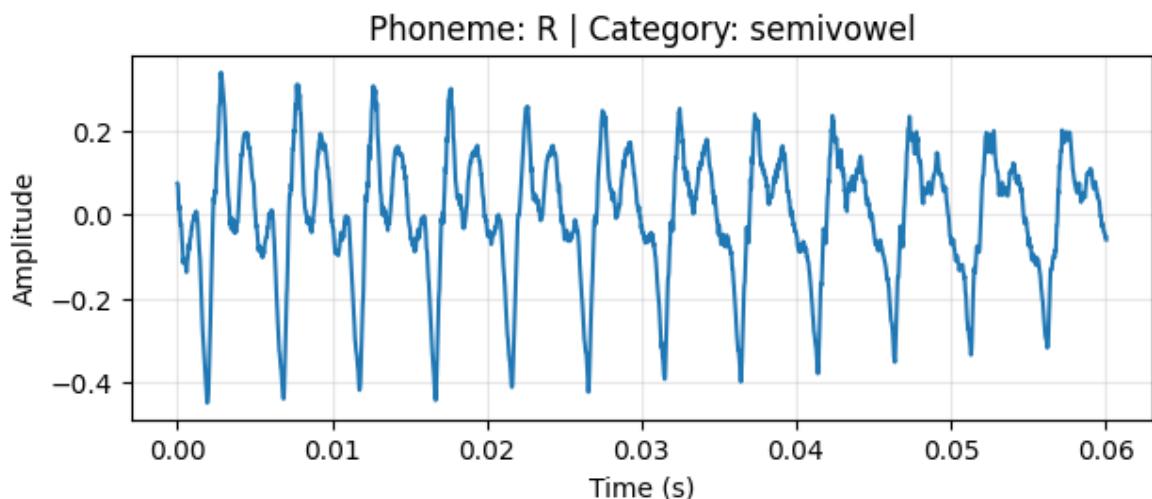
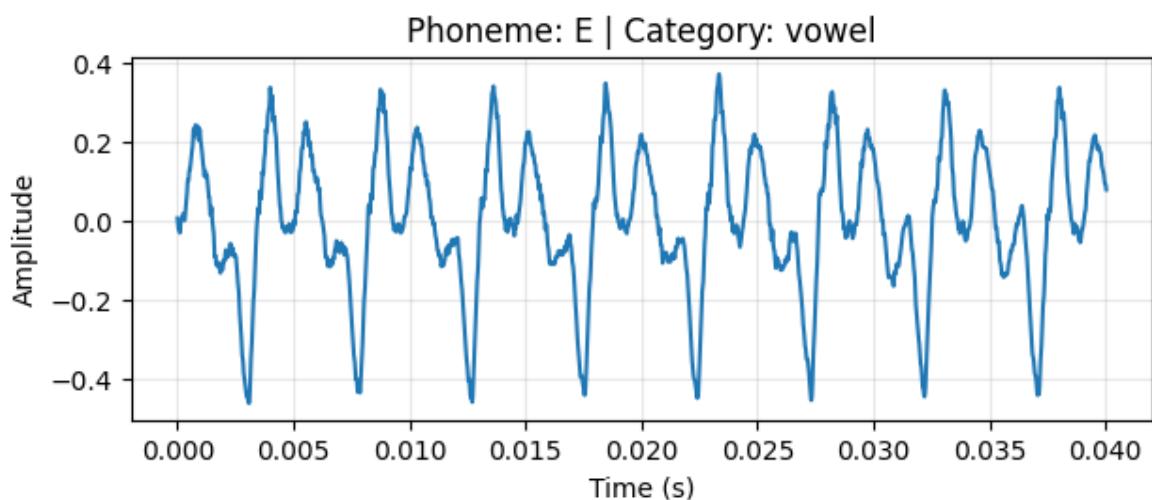
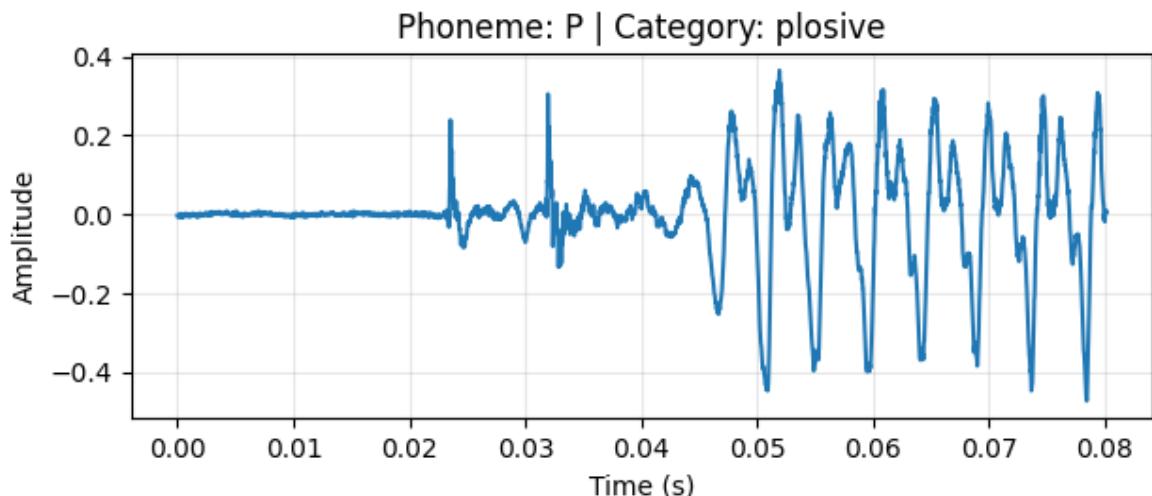


Phoneme: I | Category: vowel

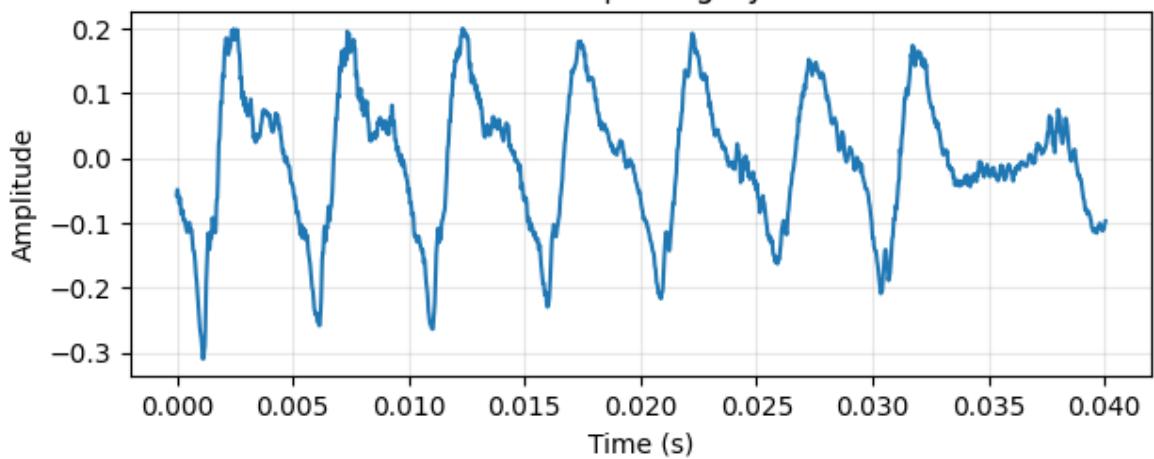


Phoneme: S | Category: fricative

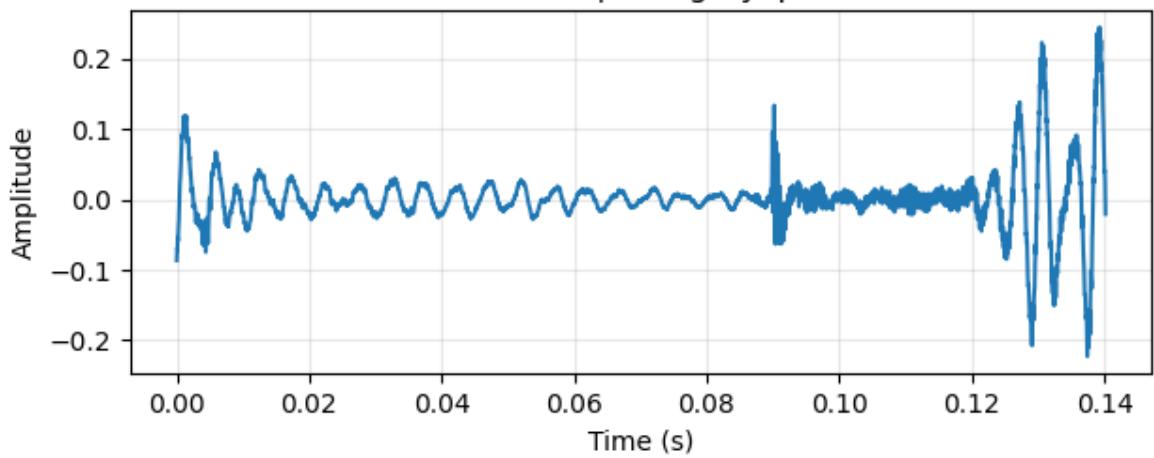




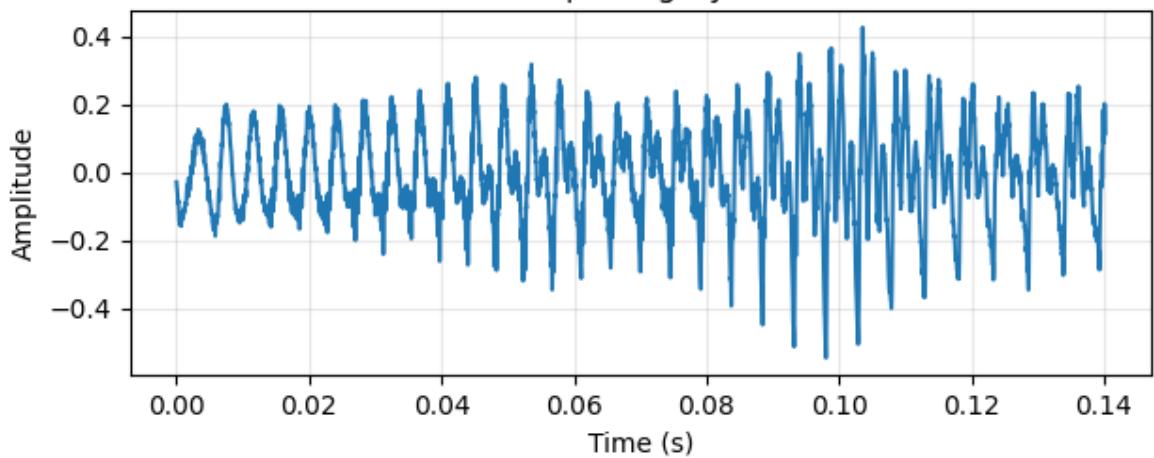
Phoneme: E | Category: vowel

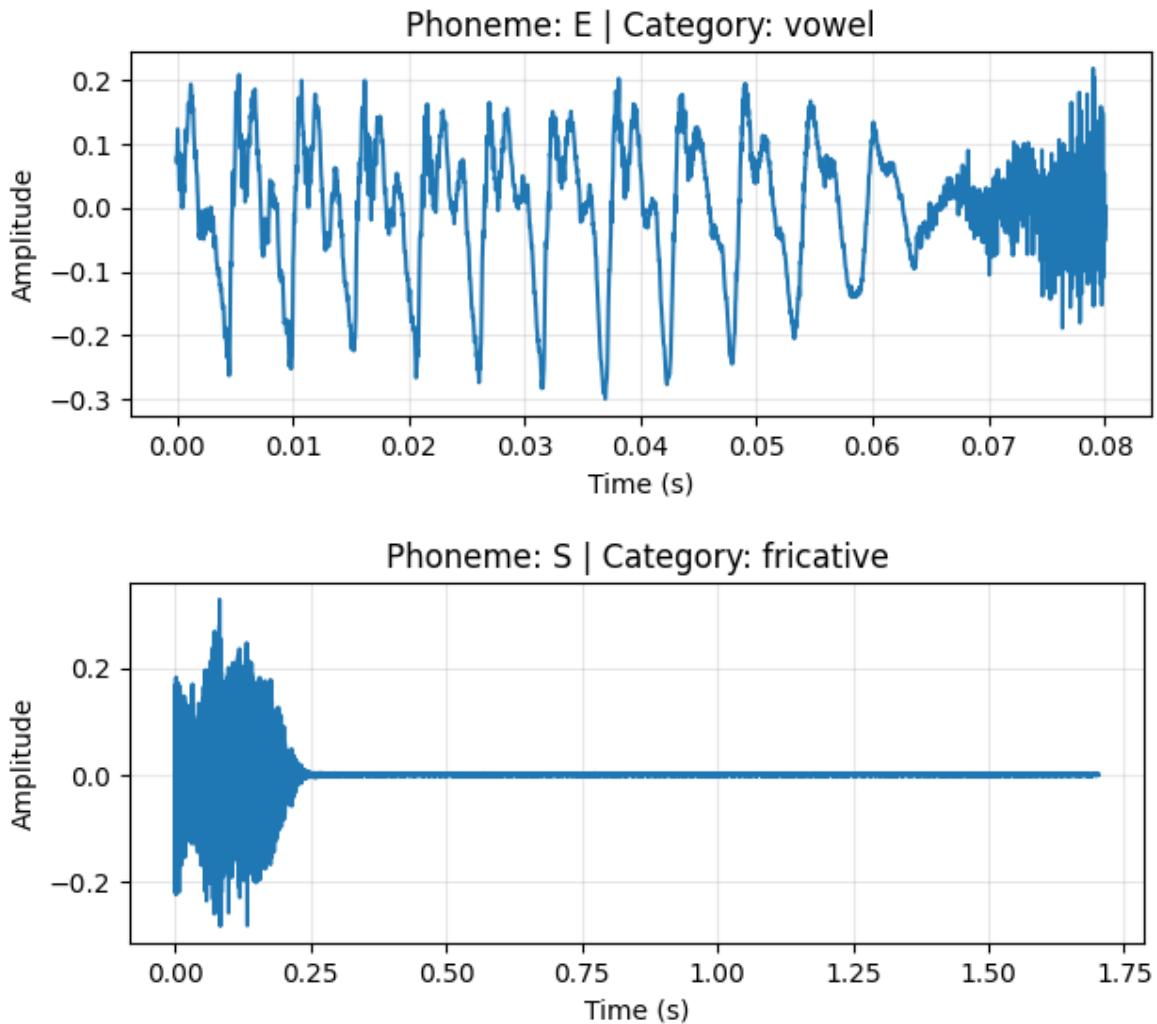


Phoneme: D | Category: plosive



Phoneme: Y | Category: semivowel





9. Voiced vs Unvoiced Comparison

```
In [9]: def get_category(p):
    vowels = ['A', 'E', 'I', 'O', 'U']
    fricatives = ['S', 'F', 'Z', 'H']
    char = p["phoneme"].upper()
    if char in vowels: return "vowel"
    if char in fricatives: return "fricative"
    return "other"

voiced = next((p for p in phoneme_intervals if get_category(p) == "vowel"), None)
unvoiced = next((p for p in phoneme_intervals if get_category(p) == "fricative"))

if voiced and unvoiced:
    print(f"Voiced segment: Phoneme '{voiced['phoneme']}' (Category: {voiced['category']})")
    plot_phoneme(
        waveform[0, voiced["start_sample"]:voiced["end_sample"]].numpy(),
        sample_rate,
        voiced["phoneme"],
        "Voiced (Vowel)"
    )

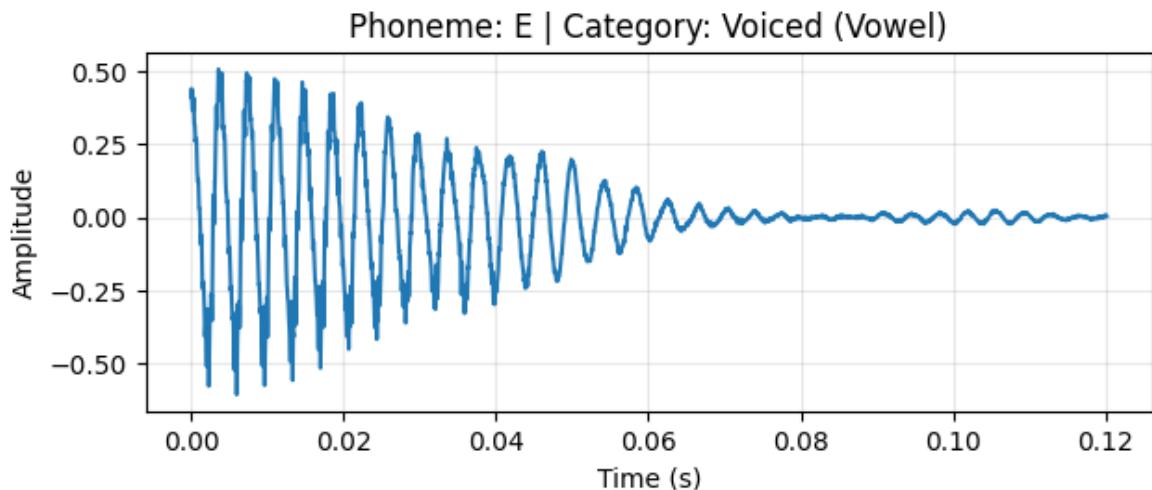
    print(f"Unvoiced segment: Phoneme '{unvoiced['phoneme']}' (Category: {unvoiced['category']})")
    plot_phoneme(
        waveform[0, unvoiced["start_sample":unvoiced["end_sample"]].numpy(),
        sample_rate,
```

```

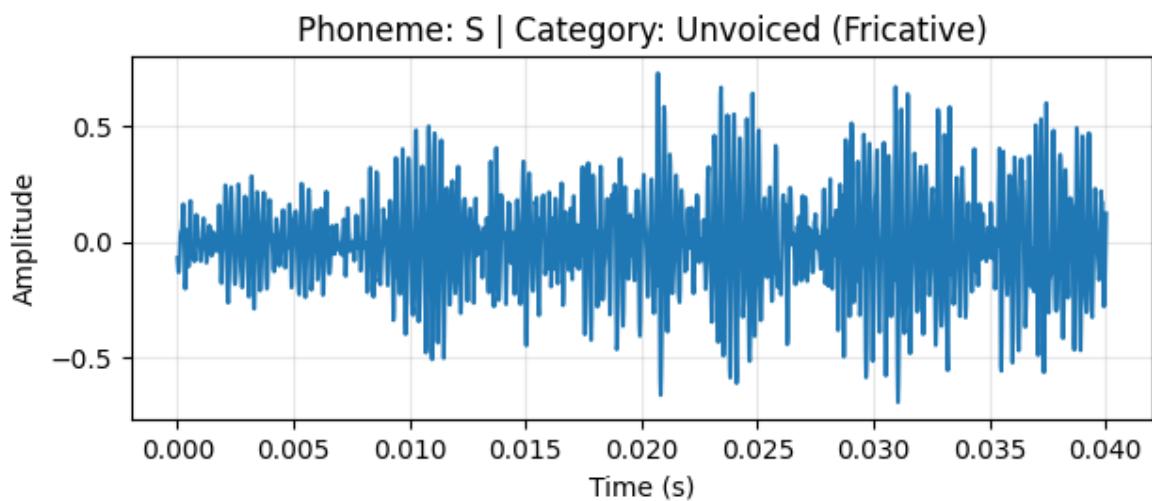
        unvoiced["phoneme"],
        "Unvoiced (Fricative)"
    )
else:
    print("Error: Could not find both a voiced and unvoiced segment.")
    print("Detected letters:", [p['phoneme'] for p in phoneme_intervals])

```

Voiced segment: Phoneme 'E' (Category: vowel)



Unvoiced segment: Phoneme 'S' (Category: fricative)



Comparison Analysis

Observed differences between voiced and unvoiced phonemes:

Periodicity:

- Voiced phonemes (vowels) show clear periodic patterns in the waveform due to regular vocal fold vibration
- Unvoiced phonemes (fricatives) display aperiodic, random fluctuations with no repeating pattern

Amplitude patterns:

- Voiced phonemes have more consistent amplitude with smooth envelope
- Unvoiced phonemes show irregular amplitude variations typical of turbulent airflow

Presence or absence of noise-like structure:

- Voiced phonemes: Absence of noise-like structure; waveform shows harmonic content
- Unvoiced phonemes: Strong noise-like structure; waveform resembles random noise from turbulent air passing through constriction

Objective 2

To analyze fricative and approximant speech sounds and infer airflow turbulence characteristics.

Record the sentence: "She sees you."

Perform the following:

- Record the given sentence at a sampling rate of 16 kHz, mono WAV format. Load the recorded speech signal.
- Identify regions corresponding to fricatives and approximants.
- Use the segmentation approach implemented in the previous lab to automatically extract phoneme-level segments.
- Isolate segments corresponding to fricatives (/s/, /ʃ/, /z/) and approximants (/j/).
- Plot the waveform of each extracted fricative and approximant segment.
- Observe qualitative differences in waveform smoothness, amplitude continuity, presence of noise-like fluctuations.
- Based on waveform observations, infer conceptually whether the airflow is turbulent (associated with higher effective Reynolds number) or Smooth/laminar (associated with lower effective Reynolds number).
- Compare fricatives and approximants in terms of waveform irregularity, periodicity, energy variation. Relate the observations to articulatory constriction differences.

In [10]:

```
import os
import librosa
import torch

audio_filename = "sentence2.wav"
# Sentence: "She sees you."

if not os.path.exists(audio_filename):
    print(f"Error: '{audio_filename}' not found. Please upload 'sentence2.wav' to the workspace")
else:
    data, sr = librosa.load(audio_filename, sr=16000)
    waveform = torch.from_numpy(data).unsqueeze(0)
    sample_rate = 16000

    waveform = waveform / torch.max(torch.abs(waveform))
    print(f"Successfully loaded {audio_filename} at {sample_rate}Hz")
```

Successfully loaded sentence2.wav at 16000Hz

11. Repeat Phoneme Recognition

```
In [11]: input_values = processor(
    waveform[0].numpy(),
    sampling_rate=sample_rate,
    return_tensors="pt"
).input_values

with torch.no_grad():
    logits = model(input_values).logits

predicted_ids = torch.argmax(logits, dim=-1)
```

```
In [12]: target_sampling_rate = 16000

if sample_rate != target_sampling_rate:
    waveform_np = waveform[0].numpy()
    waveform_resampled = librosa.resample(waveform_np, orig_sr=sample_rate, targ
    input_sr = target_sampling_rate
else:
    waveform_resampled = waveform[0].numpy()
    input_sr = sample_rate

input_values = processor(
    waveform_resampled,
    sampling_rate=input_sr,
    return_tensors="pt"
).input_values

with torch.no_grad():
    logits = model(input_values).logits

predicted_ids = torch.argmax(logits, dim=-1)
transcription = processor.decode(predicted_ids[0])

print("Recognized Phoneme Sequence:")
print(transcription)
```

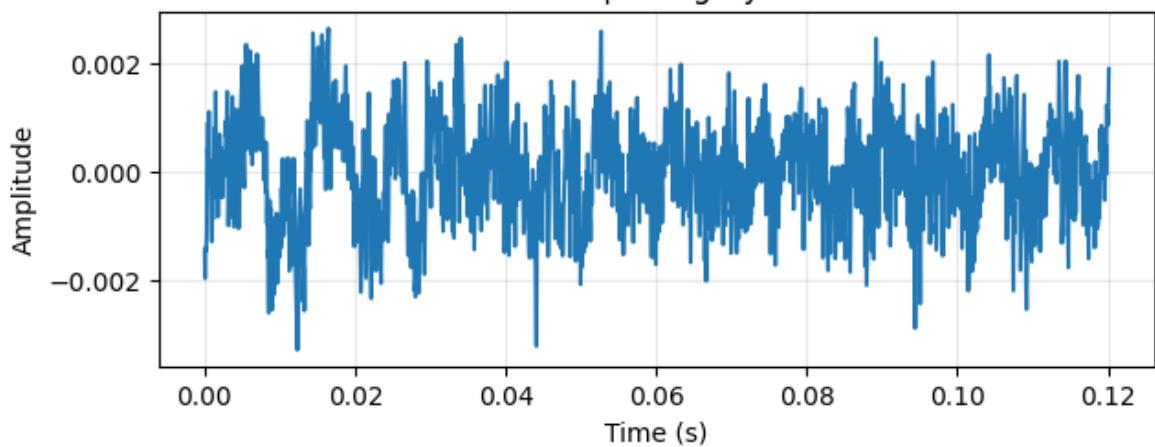
Recognized Phoneme Sequence:
SHE SEES YOU

13. Plot Fricatives & Approximants Only

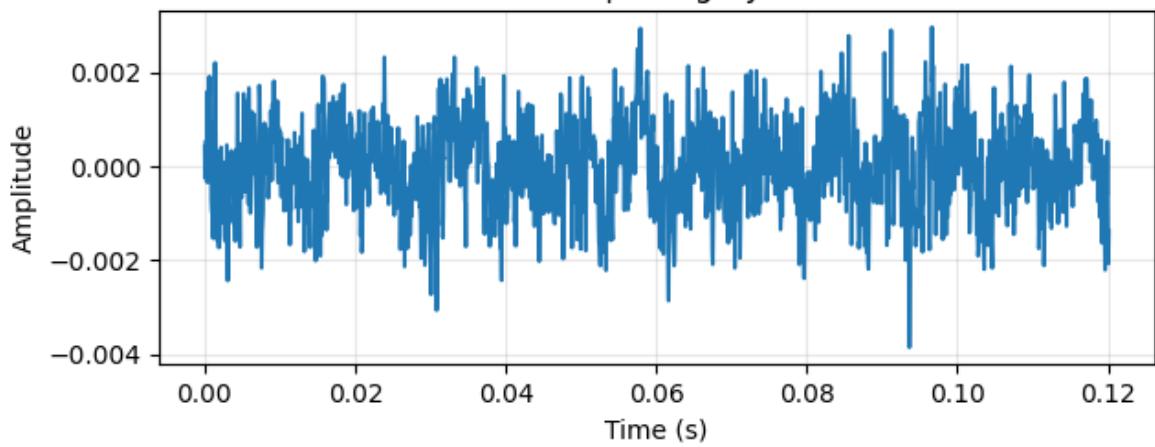
```
In [13]: for p in phoneme_intervals:
    # Check category using uppercase match
    category = phoneme_categories.get(p["phoneme"].upper(), "other")

    if category in ["fricative", "semivowel"]:
        segment = waveform[0, p["start_sample":p["end_sample"]].numpy()
        plot_phoneme(segment, sample_rate, p["phoneme"], category)
```

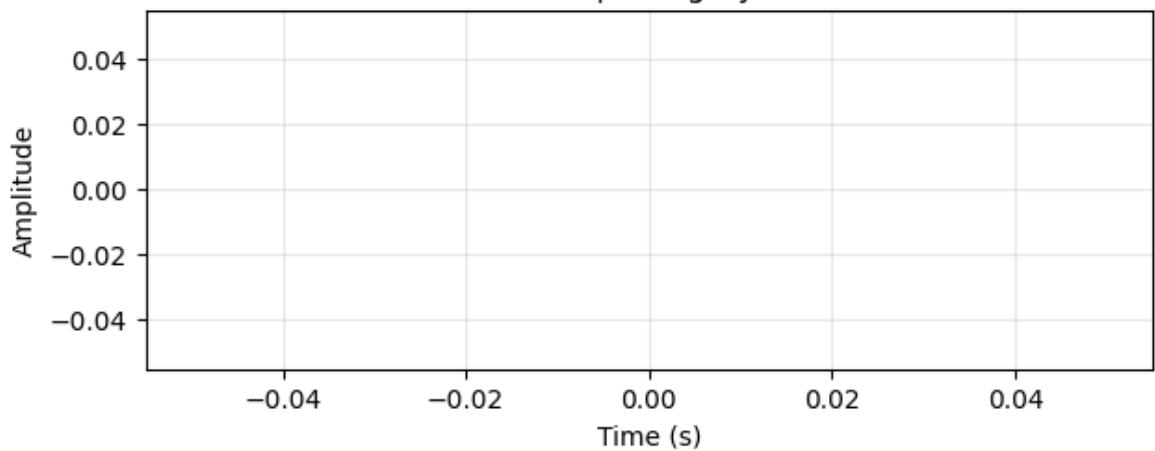
Phoneme: S | Category: fricative



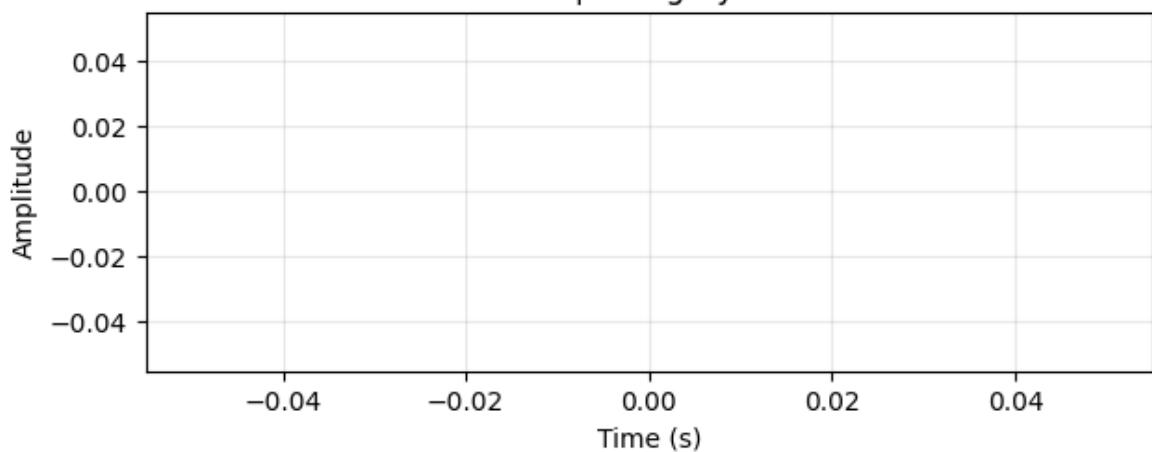
Phoneme: H | Category: fricative



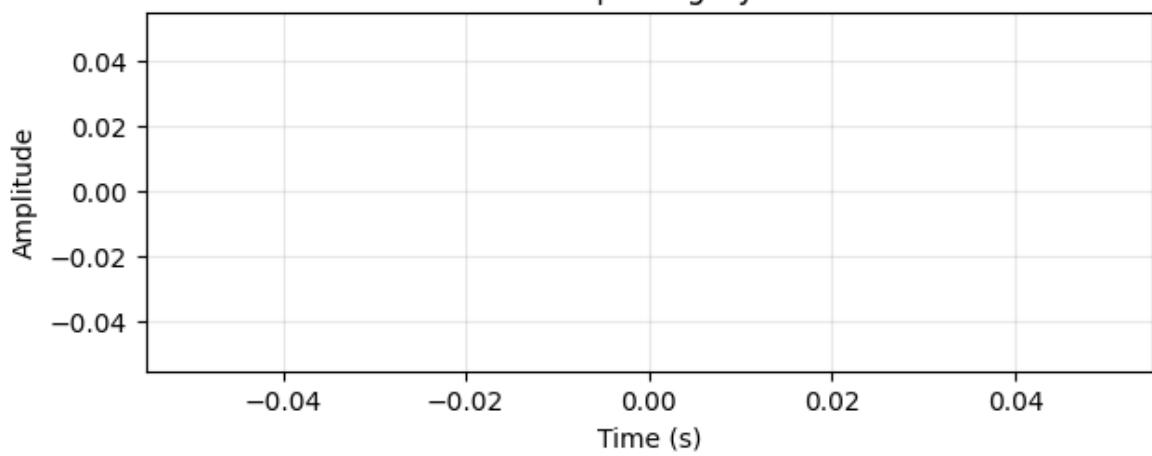
Phoneme: H | Category: fricative



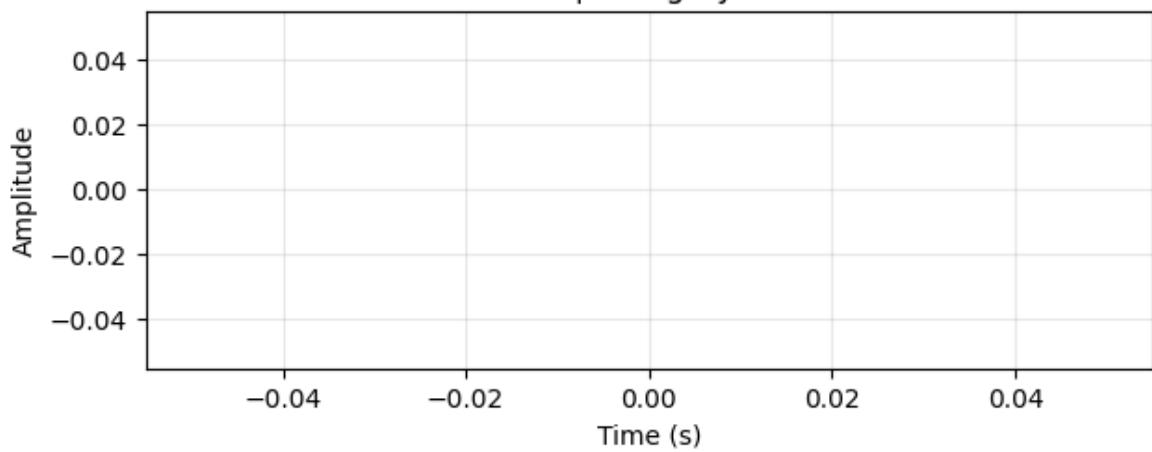
Phoneme: L | Category: semivowel



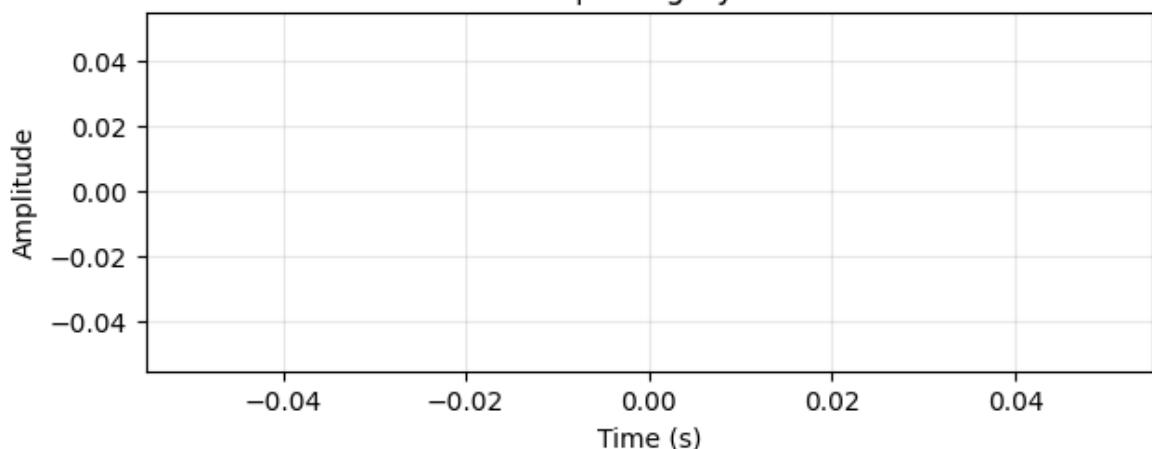
Phoneme: H | Category: fricative



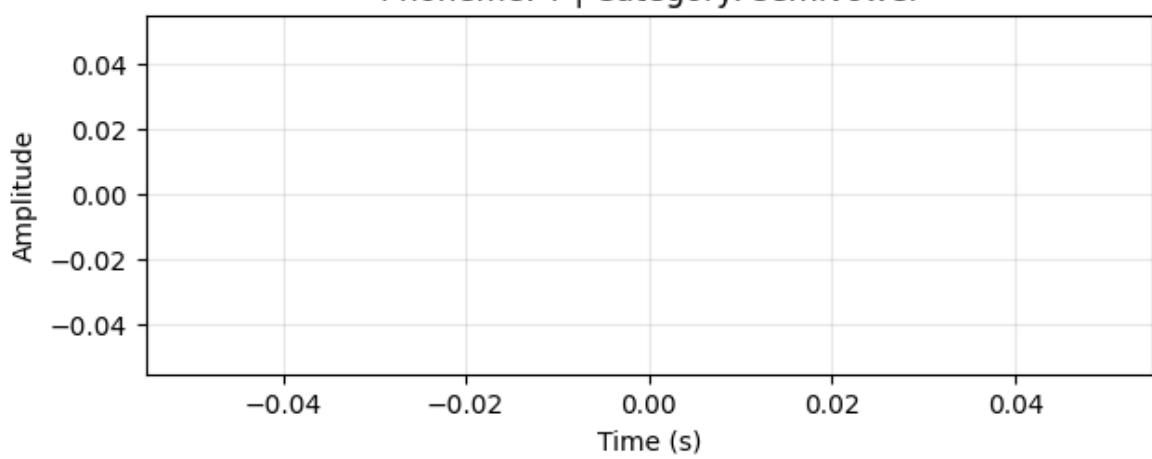
Phoneme: L | Category: semivowel



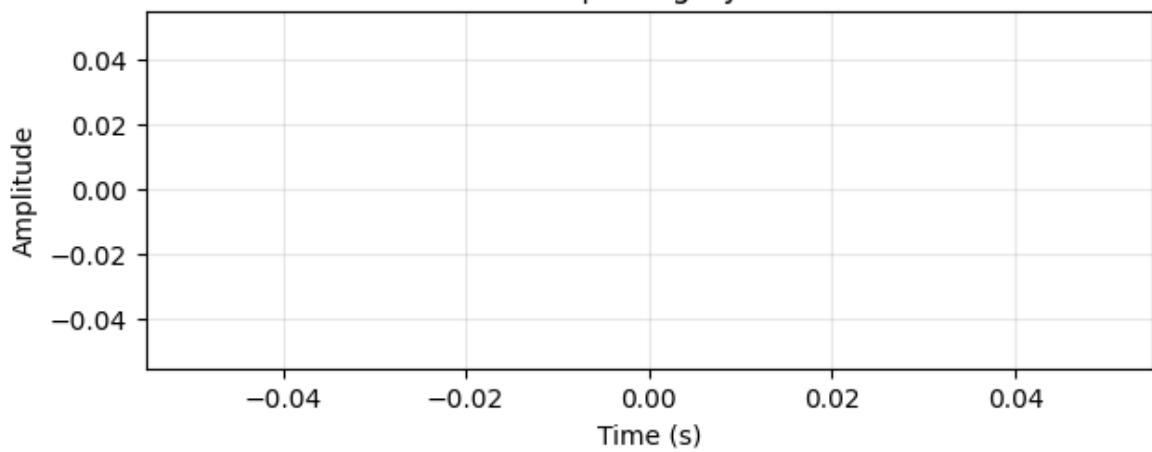
Phoneme: L | Category: semivowel



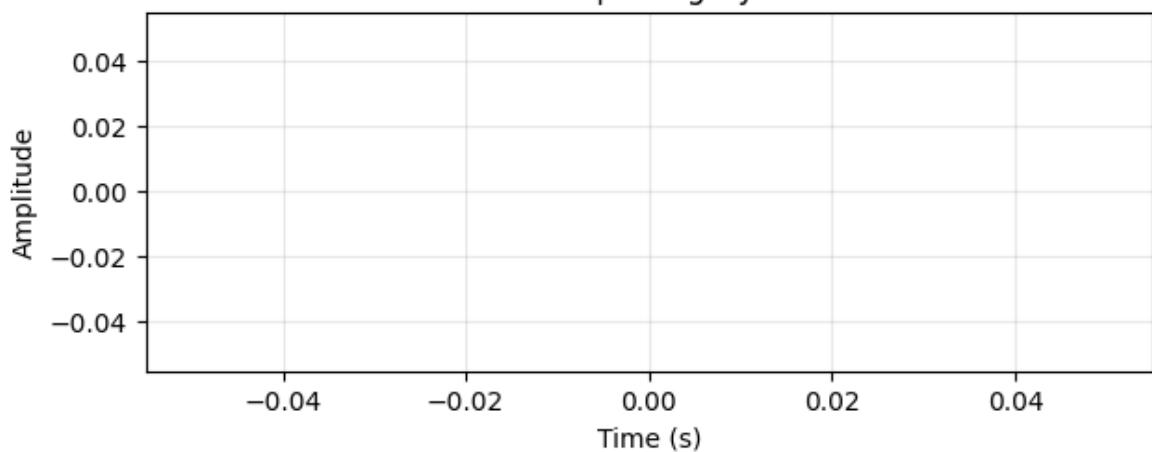
Phoneme: Y | Category: semivowel



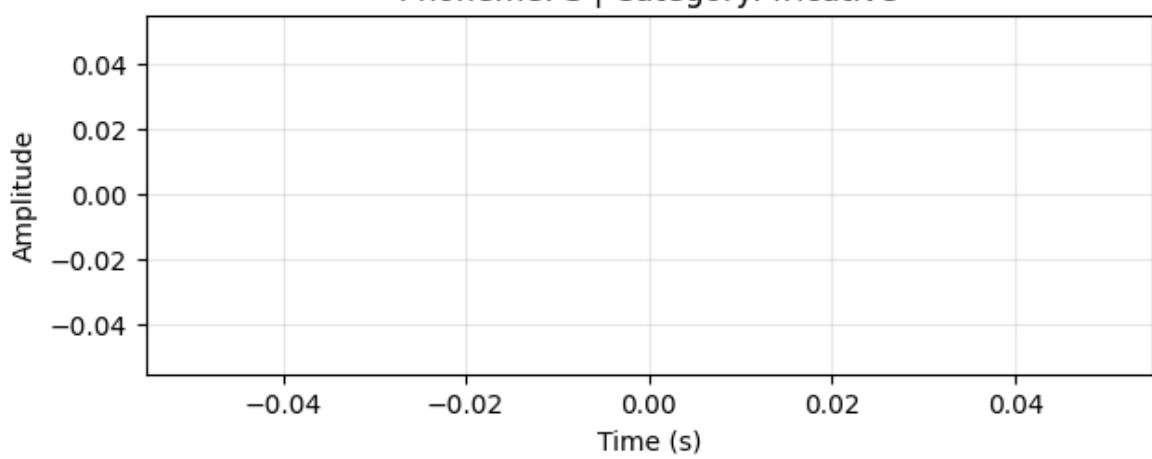
Phoneme: W | Category: semivowel



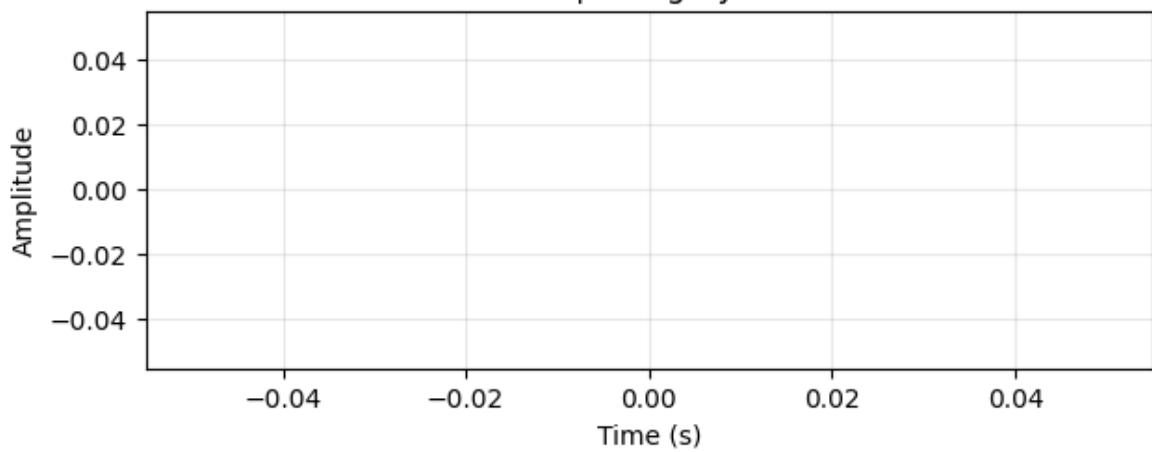
Phoneme: H | Category: fricative



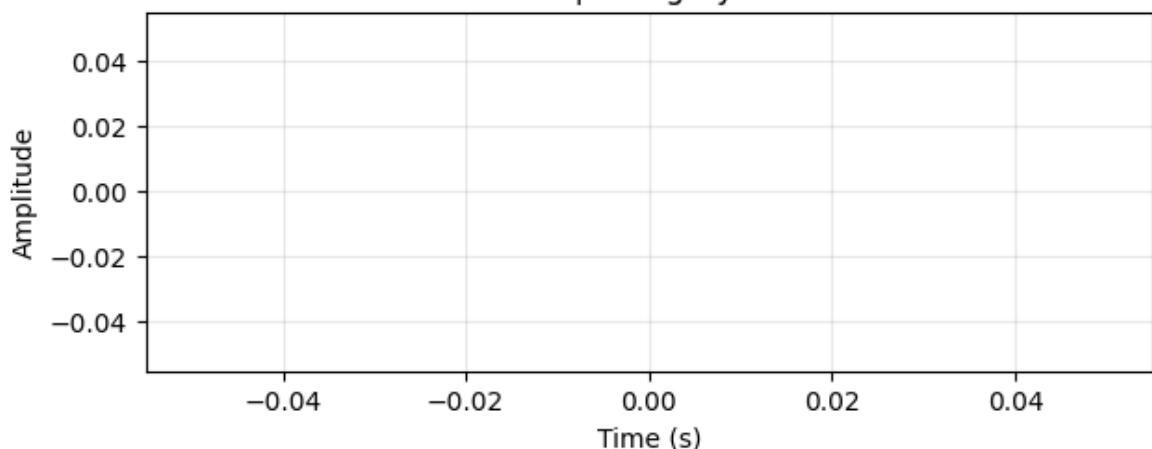
Phoneme: S | Category: fricative



Phoneme: R | Category: semivowel



Phoneme: Y | Category: semivowel



Phoneme: S | Category: fricative

