

lab4_answers

February 23, 2026

1 Speech Processing Lab Assignment - 4

1.1 Frame-wise Analysis of Speech Signal: Time-Domain Features for Voiced and Unvoiced Speech

Name: Sanjushree Rajan Reg. no.: BL.EN.U4AIE23130

Objective: Perform frame-wise analysis of a speech signal and study the behaviour of time-domain features for voiced and unvoiced speech.

1.2 Setup & Imports

```
[4]: import numpy as np
import matplotlib.pyplot as plt
import librosa
import librosa.display
import soundfile as sf
from scipy.signal.windows import hamming
import warnings
warnings.filterwarnings('ignore')

# Install if needed:
# !pip install librosa soundfile

plt.rcParams['figure.figsize'] = (14, 4)
plt.rcParams['axes.grid'] = True
print('Imports successful.')
```

Imports successful.

1.3 Load Speech Signal

You can use your own .wav file or download from LJ Speech Dataset. Place the file in the same directory and update the path below.

```
[5]: # -----
# CONFIG - update this path to point to your .wav file
AUDIO_FILE = 'speech.wav'
TARGET_SR  = 16000          # 16 kHz
```

```

DURATION = 20          # first 20 seconds
# -----
# Load (resample to 16 kHz if needed)
signal, sr = librosa.load(AUDIO_FILE, sr=TARGET_SR, duration=DURATION, mono=True)
print(f'Sample rate : {sr} Hz')
print(f'Duration : {len(signal)/sr:.2f} s')
print(f'Total samples: {len(signal)}')

```


Sample rate : 16000 Hz
Duration : 7.58 s
Total samples: 121344

1.4 Task 1 — Short-Time Processing of Speech

1.4.1 1a) Waveform Visualization (first 20 seconds)

```
[6]: time_axis = np.arange(len(signal)) / sr

plt.figure(figsize=(14, 3))
plt.plot(time_axis, signal, linewidth=0.5, color='steelblue')
plt.xlabel('Time (s)')
plt.ylabel('Amplitude')
plt.title('Speech Waveform (first 20 s)')
plt.tight_layout()
plt.show()
```



1.4.2 1b) Framing Parameters

```
[7]: frame_length_ms = 25    # ms
frame_shift_ms   = 10     # ms

frame_length = int(frame_length_ms * sr / 1000)    # samples
frame_shift   = int(frame_shift_ms   * sr / 1000)    # samples
```

```

num_frames = 1 + (len(signal) - frame_length) // frame_shift

print(f'Frame length : {frame_length} samples ({frame_length_ms} ms)')
print(f'Frame shift   : {frame_shift}   samples ({frame_shift_ms} ms)')
print(f'Total frames : {num_frames}')

```

Frame length : 400 samples (25 ms)
 Frame shift : 160 samples (10 ms)
 Total frames : 756

1.4.3 1c) Frame Segmentation

```
[8]: def segment_frames(signal, frame_length, frame_shift):
    """Return 2D array of shape (num_frames, frame_length)."""
    num_frames = 1 + (len(signal) - frame_length) // frame_shift
    frames = np.zeros((num_frames, frame_length))
    for i in range(num_frames):
        start = i * frame_shift
        frames[i] = signal[start:start + frame_length]
    return frames

frames_raw = segment_frames(signal, frame_length, frame_shift)
print(f'Frames array shape: {frames_raw.shape}')

```

Frames array shape: (756, 400)

1.4.4 1d & 1e) Window Functions + Feature Extraction Functions

```
[9]: # --- Window functions ---
hamming_win    = np.hamming(frame_length)
rect_win       = np.ones(frame_length)

# Apply windows to all frames
frames_hamming = frames_raw * hamming_win
frames_rect    = frames_raw * rect_win

# --- Feature functions ---

def compute_STE(frames):
    """Short-Time Energy: sum of squared samples per frame."""
    return np.sum(frames ** 2, axis=1)

def compute_STM(frames):
    """Short-Time Magnitude: sum of absolute values per frame."""
    return np.sum(np.abs(frames), axis=1)

def compute_ZCR(frames):
    """Zero-Crossing Rate: number of sign changes per frame."""

```

```

    signs = np.sign(frames)
    zcr = np.sum(np.abs(np.diff(signs)), axis=1) / (2 * frames.shape[1])
    return zcr

def compute_autocorrelation(frames):
    """Normalised autocorrelation at lag 1 ... frame_length-1. Returns R[1]_L
    (lag-1 value) per frame."""
    # For a compact scalar feature we use the peak value of normalised ACF
    acf_peak = []
    for frame in frames:
        acf = np.correlate(frame, frame, mode='full')
        acf = acf[len(frame)-1:] # keep non-negative lags
        if acf[0] != 0:
            acf = acf / acf[0] # normalise
        # Peak after lag 0 (excluding zero lag)
        acf_peak.append(np.max(acf[1:]) if len(acf) > 1 else 0)
    return np.array(acf_peak)

def compute_AMDF(frames):
    """Average Magnitude Difference Function - minimum value per frame (for
    feature plot)."""
    amdf_min = []
    N = frames.shape[1]
    for frame in frames:
        amdf = np.array([np.mean(np.abs(frame[k:] - frame[:N-k])) for k in
                        range(1, N//2)])
        amdf_min.append(np.min(amdf) if len(amdf) > 0 else 0)
    return np.array(amdf_min)

def compute_AMSDF(frames):
    """Average Magnitude Squared Difference Function - minimum value per frame.
    """
    amsdf_min = []
    N = frames.shape[1]
    for frame in frames:
        amsdf = np.array([np.mean((frame[k:] - frame[:N-k])**2) for k in
                        range(1, N//2)])
        amsdf_min.append(np.min(amsdf) if len(amsdf) > 0 else 0)
    return np.array(amsdf_min)

print('Feature functions defined.')

```

Feature functions defined.

```
[10]: print('Computing features for Hamming window... ')
ste_h    = compute_STE(frames_hamming)
stm_h    = compute_STM(frames_hamming)
```

```

zcr_h    = compute_ZCR(frames_hamming)
acf_h    = compute_autocorrelation(frames_hamming)
amdf_h   = compute_AMDF(frames_hamming)
amsdf_h = compute_AMSDF(frames_hamming)

print('Computing features for Hamming window...')
ste_r    = compute_STE(frames_rect)
stm_r    = compute_STM(frames_rect)
zcr_r    = compute_ZCR(frames_rect)
acf_r    = compute_autocorrelation(frames_rect)
amdf_r   = compute_AMDF(frames_rect)
amsdf_r = compute_AMSDF(frames_rect)

print('Done.')

```

Computing features for Hamming window...
 Computing features for Rectangular window...
 Done.

1.4.5 1f) Plot Frame-wise Features (Hamming vs Rectangular)

```

[11]: frame_times = np.arange(num_frames) * frame_shift_ms / 1000 # time in seconds

features = [
    ('Short-Time Energy (STE)', ste_h, ste_r),
    ('Short-Time Magnitude (STM)', stm_h, stm_r),
    ('Zero-Crossing Rate (ZCR)', zcr_h, zcr_r),
    ('ACF Peak (lag > 0)', acf_h, acf_r),
    ('AMDF Min', amdf_h, amdf_r),
    ('AMSDF Min', amsdf_h, amsdf_r),
]

fig, axes = plt.subplots(6, 2, figsize=(16, 20))
fig.suptitle('Frame-wise Time-Domain Features: Hamming (left) vs Rectangular (right)', fontsize=14, fontweight='bold')


for row, (name, feat_h, feat_r) in enumerate(features):
    axes[row, 0].plot(frame_times, feat_h, color='royalblue', linewidth=0.8)
    axes[row, 0].set_title(f'{name} - Hamming')
    axes[row, 0].set_xlabel('Time (s)')

    axes[row, 1].plot(frame_times, feat_r, color='darkorange', linewidth=0.8)
    axes[row, 1].set_title(f'{name} - Rectangular')
    axes[row, 1].set_xlabel('Time (s)')

plt.tight_layout(rect=[0, 0, 1, 0.97])
plt.show()

```

Frame-wise Time-Domain Features: Hamming (left) vs Rectangular (right)



1.4.6 1g) Analysis: Voiced vs Unvoiced Behaviour

We use a simple energy-based Voice Activity Detector (VAD) to label frames as voiced or unvoiced, then compare feature statistics.

```
[12]: # Simple energy-based VAD
ste_norm = ste_h / np.max(ste_h)
energy_threshold = 0.05 # tune if needed
voiced_mask = ste_norm >= energy_threshold
unvoiced_mask = ~voiced_mask

print(f'Voiced frames : {np.sum(voiced_mask)}')
print(f'Unvoiced frames : {np.sum(unvoiced_mask)}')

# Comparison table
print('\n--- Feature Comparison: Voiced vs Unvoiced (Hamming) ---')
for name, feat in [('STE', ste_h), ('STM', stm_h), ('ZCR', zcr_h),
                   ('ACF Peak', acf_h), ('AMDF Min', amdf_h), ('AMSDF Min', amsdf_h)]:
    v_mean = np.mean(feat[voiced_mask]) if np.any(voiced_mask) else 0
    u_mean = np.mean(feat[unvoiced_mask]) if np.any(unvoiced_mask) else 0
    print(f'{name:<10} Voiced={v_mean:.4f} Unvoiced={u_mean:.4f}')



```

Voiced frames : 256

Unvoiced frames : 500


```
--- Feature Comparison: Voiced vs Unvoiced (Hamming) ---
STE          Voiced=1.7318  Unvoiced=0.1118
STM          Voiced=16.1391  Unvoiced=3.3600
ZCR          Voiced=0.1416  Unvoiced=0.2187
ACF Peak    Voiced=0.9128  Unvoiced=0.8219
AMDF Min    Voiced=0.0141  Unvoiced=0.0035
AMSDF Min   Voiced=0.0007  Unvoiced=0.0001
```

```
[13]: # Visualise voiced/unvoiced labels against waveform + STE
fig, (ax1, ax2) = plt.subplots(2, 1, figsize=(14, 6), sharex=True)

ax1.plot(time_axis, signal, linewidth=0.5, color='gray')
ax1.set_ylabel('Amplitude')
ax1.set_title('Waveform with Voiced (blue) / Unvoiced (red) Regions')
for i, t in enumerate(frame_times):
    c = 'blue' if voiced_mask[i] else 'red'
    ax1.axvspan(t, t + frame_length_ms/1000, alpha=0.04, color=c)

ax2.plot(frame_times, ste_norm, color='steelblue', linewidth=0.8)
ax2.axhline(energy_threshold, color='red', linestyle='--', label='VAD threshold')
ax2.set_xlabel('Time (s)')
ax2.set_ylabel('Normalised STE')
ax2.set_title('Short-Time Energy (Hamming)')
ax2.legend()
```

```
plt.tight_layout()
plt.show()
```




Observations — Feature Behaviour for Voiced vs Unvoiced Speech

Feature	Voiced Speech	Unvoiced Speech	Inference
STE	High — periodic vocal cord vibration generates larger amplitude oscillations	Low — turbulent noise has small amplitude	Good voiced/unvoiced discriminator
STM	High	Low	Similar to STE; robust to squaring non-linearities
ZCR	Low — slow, quasi-periodic oscillations cross zero infrequently	High — fricative noise crosses zero rapidly	Reliable V/UV classifier; unvoiced ZCR can be 3–5× higher
ACF Peak	High (close to 1) — strong periodicity gives a prominent peak at pitch lag	Low — aperiodic signals have no dominant lag	Core feature for pitch detection
AMDF Min	Low — self-similarity at pitch period makes differences small	High — no periodic self-similarity	Computationally cheap pitch detector
AMSDF Min	Low	High	Same as AMDF but more sensitive to large differences

1.4.7 1h) Effect of Window Choice

```
[14]: # Show side-by-side STE for one voiced segment
v_frames = np.where(voiced_mask)[0]
if len(v_frames) > 0:
    start_frame = v_frames[0]
    end_frame   = min(v_frames[0] + 100, num_frames)
    seg_times   = frame_times[start_frame:end_frame]

    plt.figure(figsize=(14, 4))
    plt.plot(seg_times, ste_h[start_frame:end_frame],   label='Hamming', linewidth=1.2)
    plt.plot(seg_times, ste_r[start_frame:end_frame],   label='Rectangular', linewidth=1.2, linestyle='--')
    plt.xlabel('Time (s)')
    plt.ylabel('STE')
    plt.title('STE Comparison: Hamming vs Rectangular (voiced segment)')
    plt.legend()
    plt.tight_layout()
    plt.show()
else:
    print('No voiced frames detected - lower the VAD threshold.')
```



Window Effect Discussion

- **Hamming window** tapers to near-zero at frame edges, reducing spectral leakage and producing **smoother feature contours** with better discrimination between voiced and unvoiced regions.
- **Rectangular window** treats all samples equally. The abrupt edges cause **Gibbs-like ringing** and artificial discontinuities at frame boundaries, making feature curves slightly noisier and less smooth.
- For pitch-related features (ACF, AMDF, AMSDF), the Hamming window suppresses spurious peaks caused by edge effects, resulting in more reliable pitch period estimation.
- For ZCR and energy, the rectangular window may appear to give slightly higher values due to edge artifacts, but discrimination capability remains similar.

- **Conclusion:** Hamming window is preferred for speech analysis due to better spectral containment and smoother feature trajectories.
-

1.5 Task 2 — Periodicity Analysis

1.5.1 2a) Identify Periodic Frames


```
[15]: ACF_THRESHOLD = 0.3    # frames with ACF peak > this are considered periodic

periodic_mask = acf_h >= ACF_THRESHOLD
print(f'Periodic (voiced) frames : {np.sum(periodic_mask)}')
print(f'Aperiodic (unvoiced) frames: {np.sum(~periodic_mask)}')

plt.figure(figsize=(14, 3))
plt.plot(frame_times, acf_h, linewidth=0.8, label='ACF Peak')
plt.axhline(ACF_THRESHOLD, color='red', linestyle='--', label=f'Threshold_{ACF_THRESHOLD}')
plt.fill_between(frame_times, acf_h, where=periodic_mask, alpha=0.3, color='green', label='Periodic')
plt.xlabel('Time (s)')
plt.ylabel('Normalised ACF Peak')
plt.title('Periodicity Detection via ACF Peak')
plt.legend()
plt.tight_layout()
plt.show()
```

Periodic (voiced) frames : 731

Aperiodic (unvoiced) frames: 25



1.5.2 2b) Periodicity in ACF, AMDF, AMSDF — Single-Frame Illustration

```
[16]: # Pick one clearly periodic frame and one clearly aperiodic frame
if np.any(periodic_mask) and np.any(~periodic_mask):
    voiced_idx = np.where(periodic_mask)[0][len(np.where(periodic_mask)[0])//2] # middle voiced
    unvoiced_idx = np.where(~periodic_mask)[0][0]
```

```

frame_v = frames_hamming[voiced_idx]
frame_u = frames_hamming[unvoiced_idx]

N = frame_length
lags = np.arange(0, N)

def full_acf(frame):
    a = np.correlate(frame, frame, mode='full')
    a = a[len(frame)-1:]
    return a / a[0] if a[0] != 0 else a

def full_amdf(frame):
    return np.array([np.mean(np.abs(frame[k:] - frame[:N-k])) for k in range(1, N//2)])

def full_amsdf(frame):
    return np.array([np.mean((frame[k:] - frame[:N-k])**2) for k in range(1, N//2)])

acf_v, acf_u = full_acf(frame_v), full_acf(frame_u)
amdf_v, amdf_u = full_amdf(frame_v), full_amdf(frame_u)
amsdf_v, amsdf_u = full_amsdf(frame_v), full_amsdf(frame_u)

lag_axis = np.arange(len(acf_v))
diff_axis = np.arange(1, N//2)

fig, axes = plt.subplots(3, 2, figsize=(14, 10))
fig.suptitle('Periodicity Functions: Voiced (left) vs Unvoiced (right)', fontsize=13, fontweight='bold')

axes[0,0].plot(lag_axis, acf_v, 'royalblue'); axes[0,0].set_title('ACF - Voiced')
axes[0,1].plot(lag_axis, acf_u, 'darkorange'); axes[0,1].set_title('ACF - Unvoiced')
axes[1,0].plot(diff_axis, amdf_v, 'royalblue'); axes[1,0].set_title('AMDF - Voiced')
axes[1,1].plot(diff_axis, amdf_u, 'darkorange'); axes[1,1].set_title('AMDF - Unvoiced')
axes[2,0].plot(diff_axis, amsdf_v, 'royalblue'); axes[2,0].set_title('AMSDF - Voiced')
axes[2,1].plot(diff_axis, amsdf_u, 'darkorange'); axes[2,1].set_title('AMSDF - Unvoiced')


for ax in axes.flat:
    ax.set_xlabel('Lag (samples)')
plt.tight_layout(rect=[0, 0, 1, 0.96])

```

```

    plt.show()
else:
    print('Adjust thresholds to identify voiced/unvoiced frames.')

```



1.5.3 2c) Distinguishing Periodic vs Aperiodic Frames

Measure	Periodic (Voiced)	Aperiodic (Unvoiced)
ACF	Clear secondary peaks at multiples of pitch period T	Rapidly decaying, noisy, no dominant peak
AMDF	Deep valley (minimum) at lag = T	Roughly flat — no significant minimum
AMSDF	Sharp minimum at lag = T	Flat and elevated throughout

Explanation: In voiced speech the glottal source is quasi-periodic with period T (pitch period). When the analysis window aligns at lag T, the signal is correlated with itself, giving ACF a peak and AMDF/AMSDF a minimum. Unvoiced speech is spectrally similar to white noise — no such self-similarity exists at any lag.

1.6 Task 3 — Pitch Estimation

1.6.1 3a & 3b) Pitch Period and Frequency Estimation

```
[17]: def estimate_pitch_acf(frame, sr, f0_min=60, f0_max=400):
    """Estimate pitch period (samples) and frequency (Hz) using ACF."""
    N = len(frame)
    acf = np.correlate(frame, frame, mode='full')
    acf = acf[N-1:]          # non-negative lags
    if acf[0] != 0:
        acf = acf / acf[0]

    lag_min = int(sr / f0_max)    # shortest allowable pitch period
    lag_max = int(sr / f0_min)    # longest allowable pitch period
    lag_max = min(lag_max, N-1)

    if lag_min >= lag_max:
        return None, None

    search = acf[lag_min:lag_max+1]
    pitch_lag = np.argmax(search) + lag_min
    pitch_freq = sr / pitch_lag if pitch_lag > 0 else 0
    return pitch_lag, pitch_freq

def estimate_pitch_amdf(frame, sr, f0_min=60, f0_max=400):
    """Estimate pitch using first minimum of AMDF."""
    N = len(frame)
    lag_min = int(sr / f0_max)
    lag_max = int(sr / f0_min)
    lag_max = min(lag_max, N//2 - 1)
    if lag_min >= lag_max:
        return None, None
    amdf = np.array([np.mean(np.abs(frame[k:] - frame[:N-k])) for k in
                    range(lag_min, lag_max+1)])
    pitch_lag = np.argmin(amdf) + lag_min
    pitch_freq = sr / pitch_lag if pitch_lag > 0 else 0
    return pitch_lag, pitch_freq

# Estimate pitch for all periodic frames
pitch_freq_acf = np.zeros(num_frames)
pitch_freq_amdf = np.zeros(num_frames)

for i in range(num_frames):
    if periodic_mask[i]:
        frame = frames_hamming[i]
        _, f_acf = estimate_pitch_acf(frame, sr)
        _, f_amdf = estimate_pitch_amdf(frame, sr)
        pitch_freq_acf[i] = f_acf if f_acf else 0
```

```


pitch_freq_amdf[i] = f_amdf if f_amdf else 0

print('Pitch estimation complete.')
valid_acf = pitch_freq_acf[pitch_freq_acf > 0]
valid_amdf = pitch_freq_amdf[pitch_freq_amdf > 0]
if len(valid_acf) > 0: print(f'ACF - Mean F0: {np.mean(valid_acf):.1f} Hz | '
    'Median: {np.median(valid_acf):.1f} Hz')
if len(valid_amdf) > 0: print(f'AMDF - Mean F0: {np.mean(valid_amdf):.1f} Hz | '
    'Median: {np.median(valid_amdf):.1f} Hz')

```

Pitch estimation complete.
ACF - Mean F0: 216.9 Hz | Median: 202.5 Hz
AMDF - Mean F0: 228.0 Hz | Median: 213.3 Hz

```
[18]: plt.figure(figsize=(14, 4))
plt.plot(frame_times, pitch_freq_acf, label='Pitch (ACF)', linewidth=0.9, color='royalblue')
plt.plot(frame_times, pitch_freq_amdf, label='Pitch (AMDF)', linewidth=0.9, color='tomato', linestyle='--')
plt.xlabel('Time (s)')
plt.ylabel('Fundamental Frequency F0 (Hz)')
plt.title('Pitch Track - ACF vs AMDF')
plt.ylim(0, 500)
plt.legend()
plt.tight_layout()
plt.show()
```



Interpretation

- Typical **male pitch**: 85–180 Hz. Typical **female pitch**: 165–255 Hz.
- Frames with estimated F0 in the physiologically plausible range (60–400 Hz) are reliable voiced frames.
- ACF and AMDF generally agree on the pitch contour; differences arise at frame transitions where one method may detect octave errors (e.g., picking a harmonic instead of the fundamental).

1.7 Discussion

1.7.1 a) Voiced vs Unvoiced Speech: Energy, ZCR, and Periodicity

Voiced speech is produced by quasi-periodic vocal cord vibrations exciting the vocal tract. This results in: - **High STE/STM** due to concentrated energy in harmonics. - **Low ZCR** because the waveform oscillates slowly at the pitch frequency (e.g., 100–300 zero-crossings per second vs 2000–5000 for unvoiced). - **High periodicity** reflected as prominent ACF peaks and sharp AMDF/AMSDF minima at the pitch lag.

Unvoiced speech (fricatives, stops) is generated by turbulent airflow and shows: - **Low STE/STM** — broadband noise has smaller amplitude. - **High ZCR** — rapid random oscillations cross zero frequently. - **No periodicity** — ACF decays immediately; AMDF/AMSDF remain flat.

1.7.2 b) Periodicity-Based Features for Pitch Detection

ACF, AMDF, and AMSDF exploit the temporal self-similarity of voiced speech. A periodic signal with period T produces: - An ACF peak at lag T (and its multiples). - An AMDF/AMSDF minimum at lag T — because the frame subtracted from itself shifted by exactly one period yields near-zero difference.

This makes the lag at which the peak/minimum occurs a direct measure of the pitch period: $F_0 = F_s / T$.

1.7.3 c) Reliability of Time-Domain Features for Pitch Detection

Feature	Reliability	Notes
STE/STM	Low for pitch	Good for V/UV detection, not for pitch period
ZCR	Low for pitch	Useful for V/UV, but pitch from ZCR is unreliable
ACF	High	Gold standard for time-domain pitch; susceptible to octave errors
AMDF	High	Computationally cheaper than ACF; sometimes preferred in real-time systems
AMSDF	High	More sensitive than AMDF; can be noisier in low SNR conditions

Conclusion: ACF is the most widely used and reliable time-domain pitch estimator. AMDF is a computationally efficient alternative. Combined use of multiple features (e.g., ACF for detection

+ AMDF for confirmation) improves robustness. All time-domain methods are sensitive to noise; for noisy conditions, frequency-domain methods (CEPSTRUM, HPS) or deep learning approaches are preferred.